

UNIVERSITY OF  
PORTSMOUTH

# AI 101

Becky Canning

Institute of Cosmology and Gravitation



© Becky Canning

# With huge thanks to:

- Ben, Coleman, Lorenza, Obinna, Xan

# AI 101

- Key AI things to know:
  - Concepts of: feature engineering, anatomy of a model, and regularisation
- For EO:
  - Querying web APIs with Python
  - Identifying things in images
  - What might we use for time series?
  - The pitfalls often found with EO data
- For space:
  - TinyML models

# AI 101

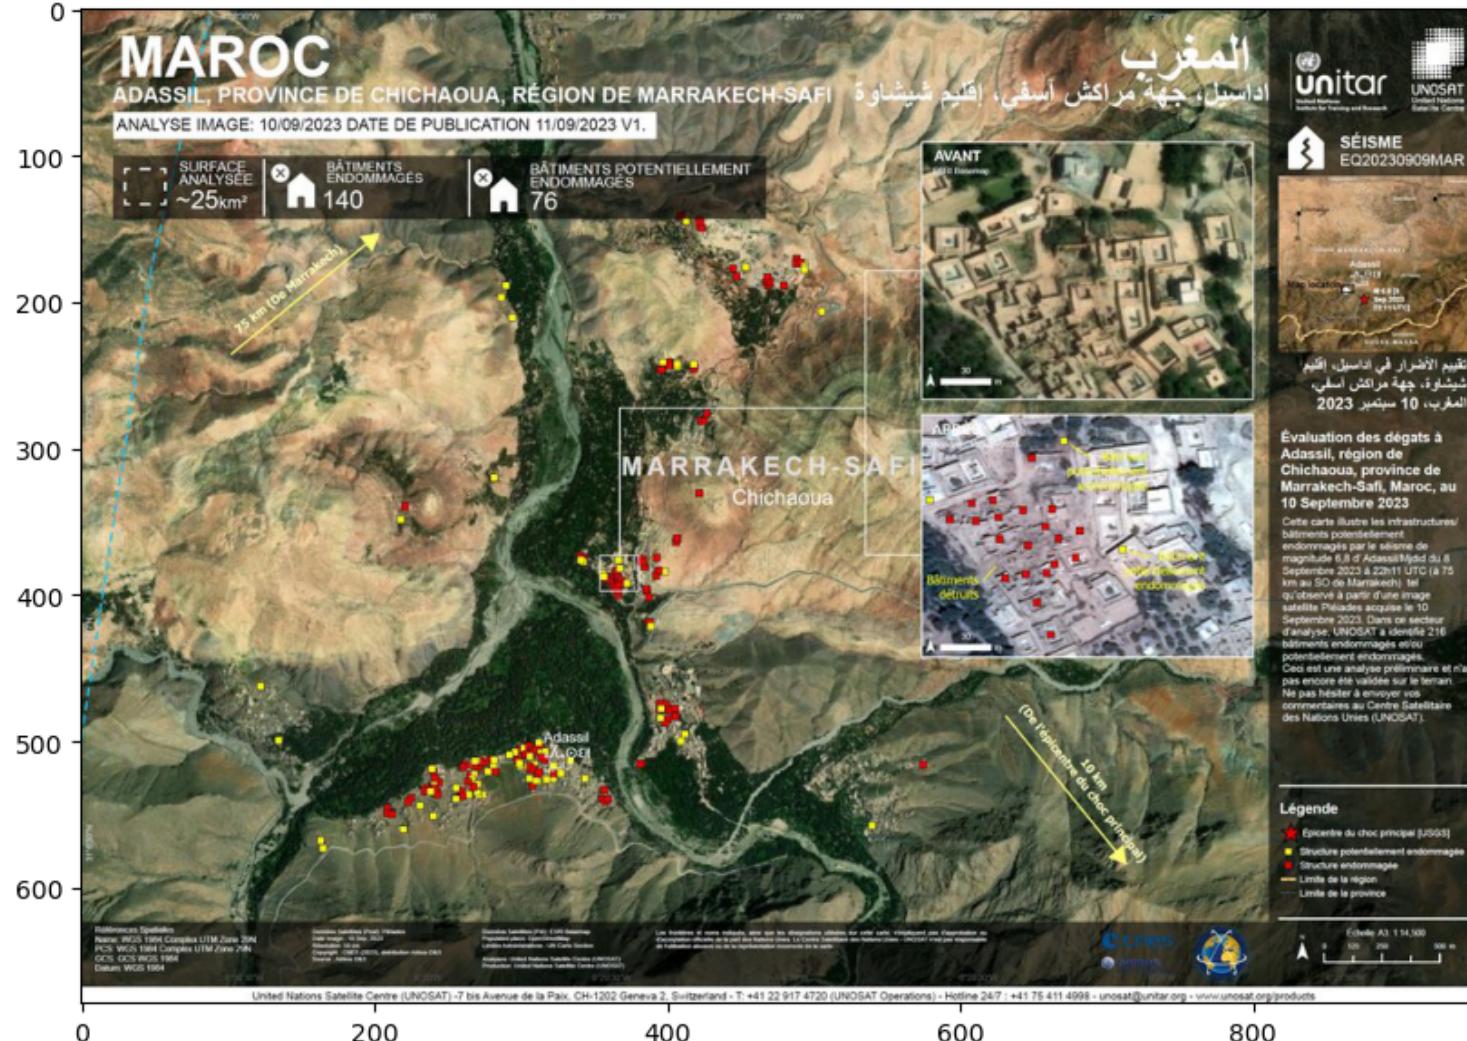
- Key AI things to know:
  - Concepts of: feature engineering, anatomy of a model, and regularisation
- For EO:
  - Querying web APIs with Python
  - Identifying things in images
  - What might we use for time series?
  - The pitfalls often found with EO data
- For space:
  - TinyML models

- Clustering
- Classification

- DL case study with thanks to Obinna

# Starting from the end

- A DL model for Earth Observation data



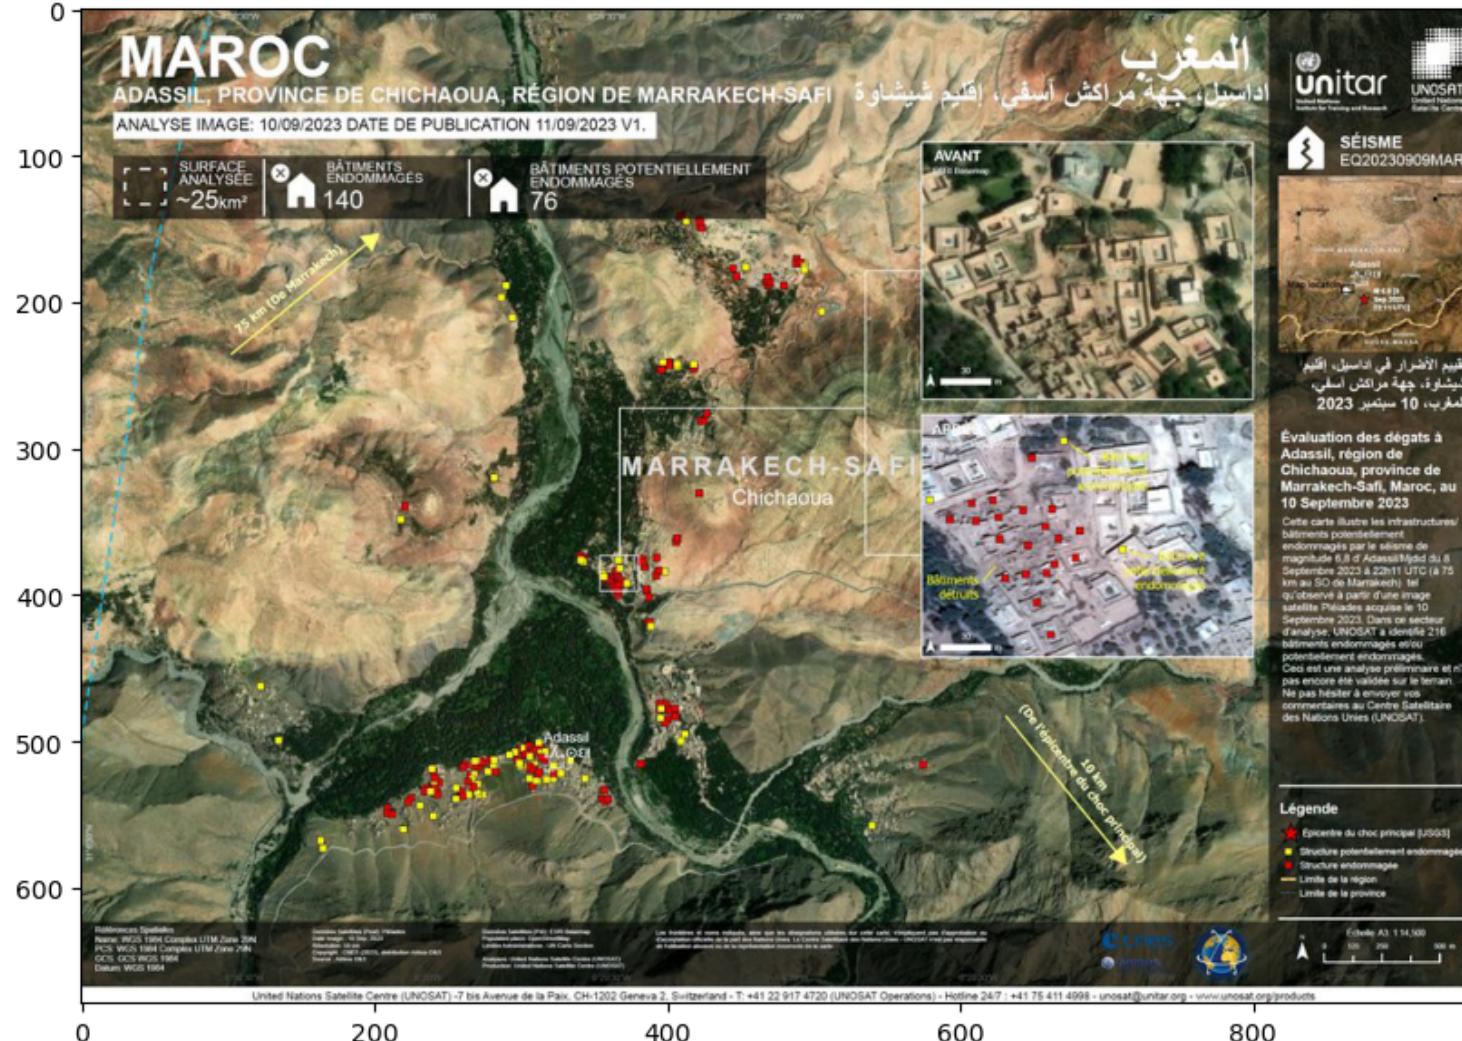
There was a 6.8 magnitude earthquake that struck the Atlas Mountains, about 75 km from Marrakech, Morocco on 8 September, 2023

Can we use satellite DL to better understand the destruction to the buildings. Identification of secure and insecure structures could help people affected?

- DL case study with thanks to Obinna

# Starting from the end

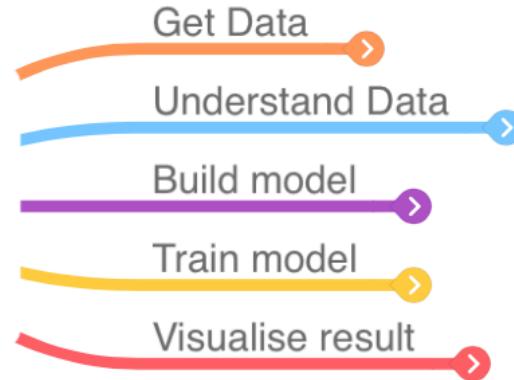
- A DL model for Earth Observation data



What are the steps we would take?

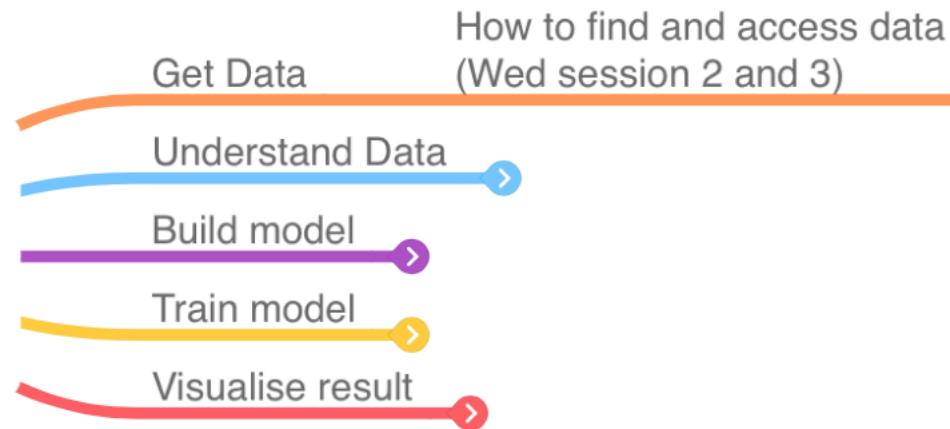
# Starting from the end

DL Case Study - what is  
the problem under study  
(classification, regression,  
clustering, ...)



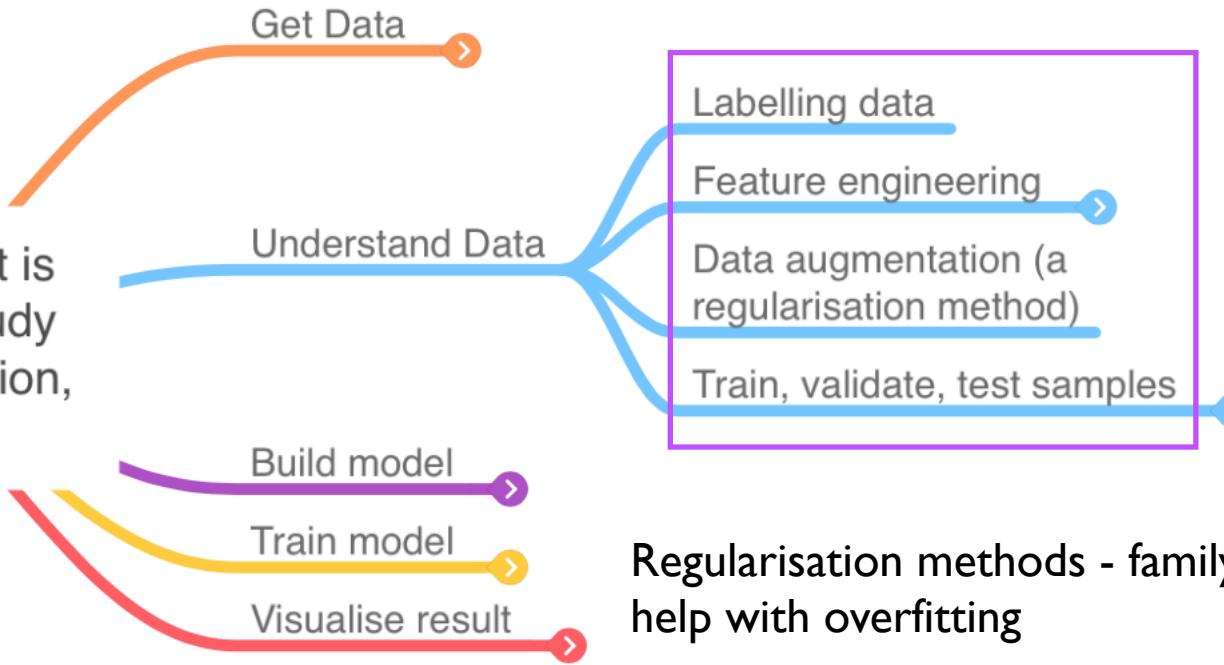
# Starting from the end

DL Case Study - what is  
the problem under study  
(classification, regression,  
clustering, ...)



# Starting from the end

DL Case Study - what is the problem under study (classification, regression, clustering, ...)

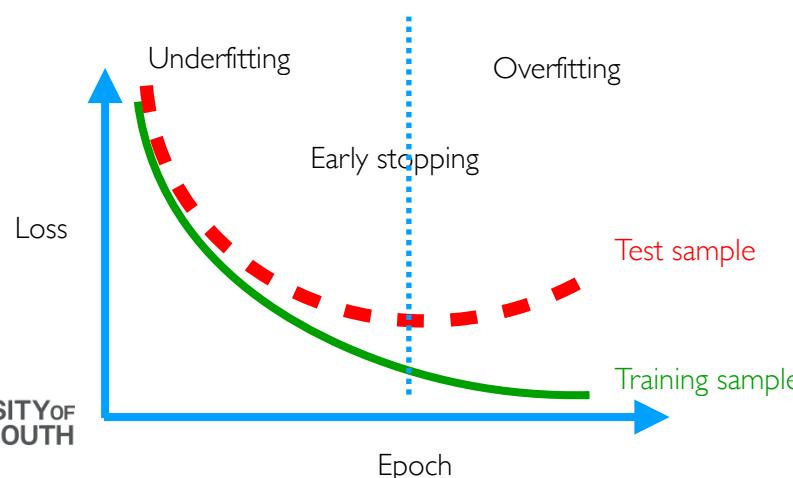


**Regularisation methods** - family of different methods which help with overfitting

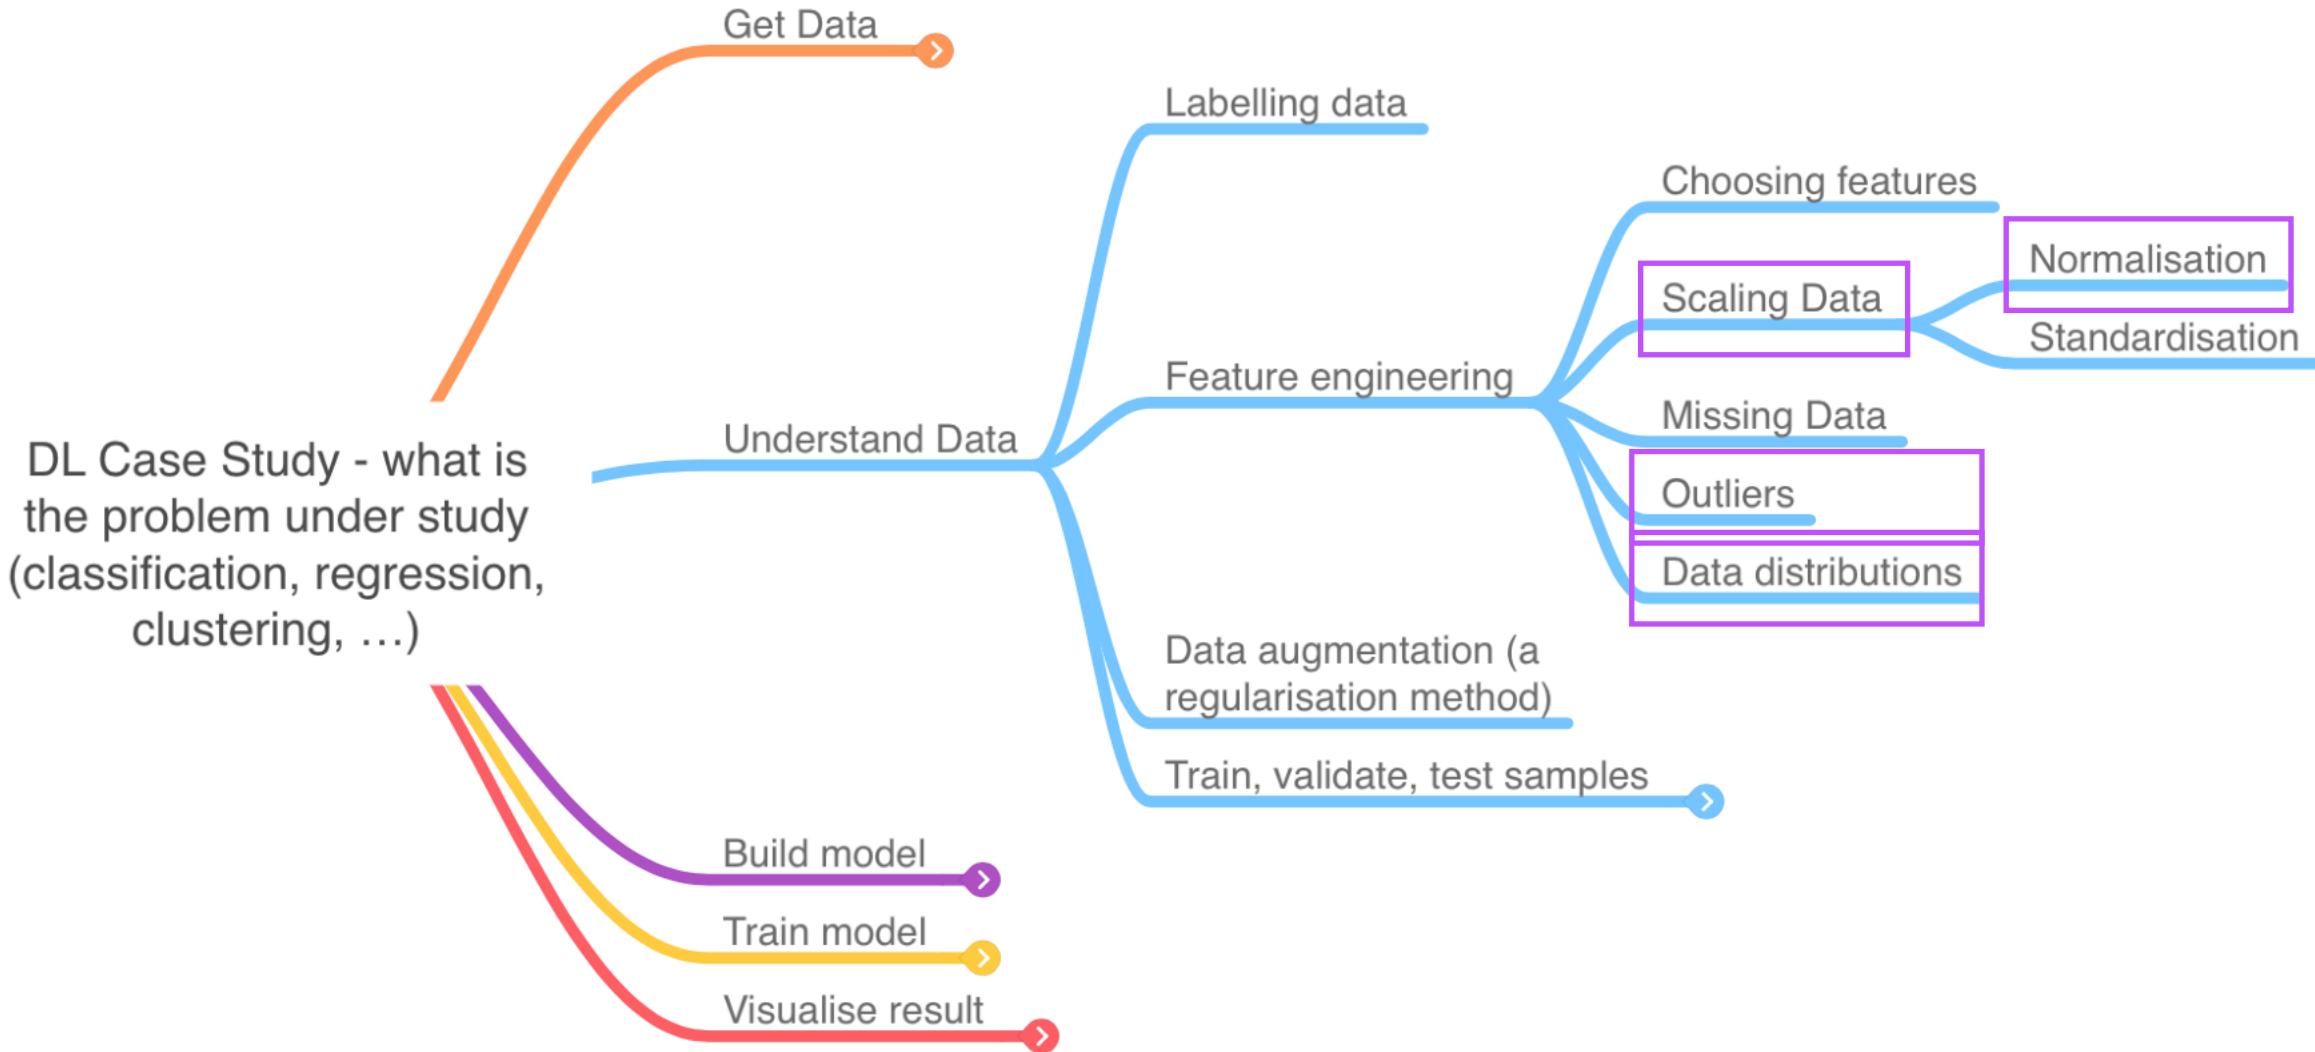
Trade off between bias and variance

**Bias**: high bias is a high training error - accuracy on training set is low

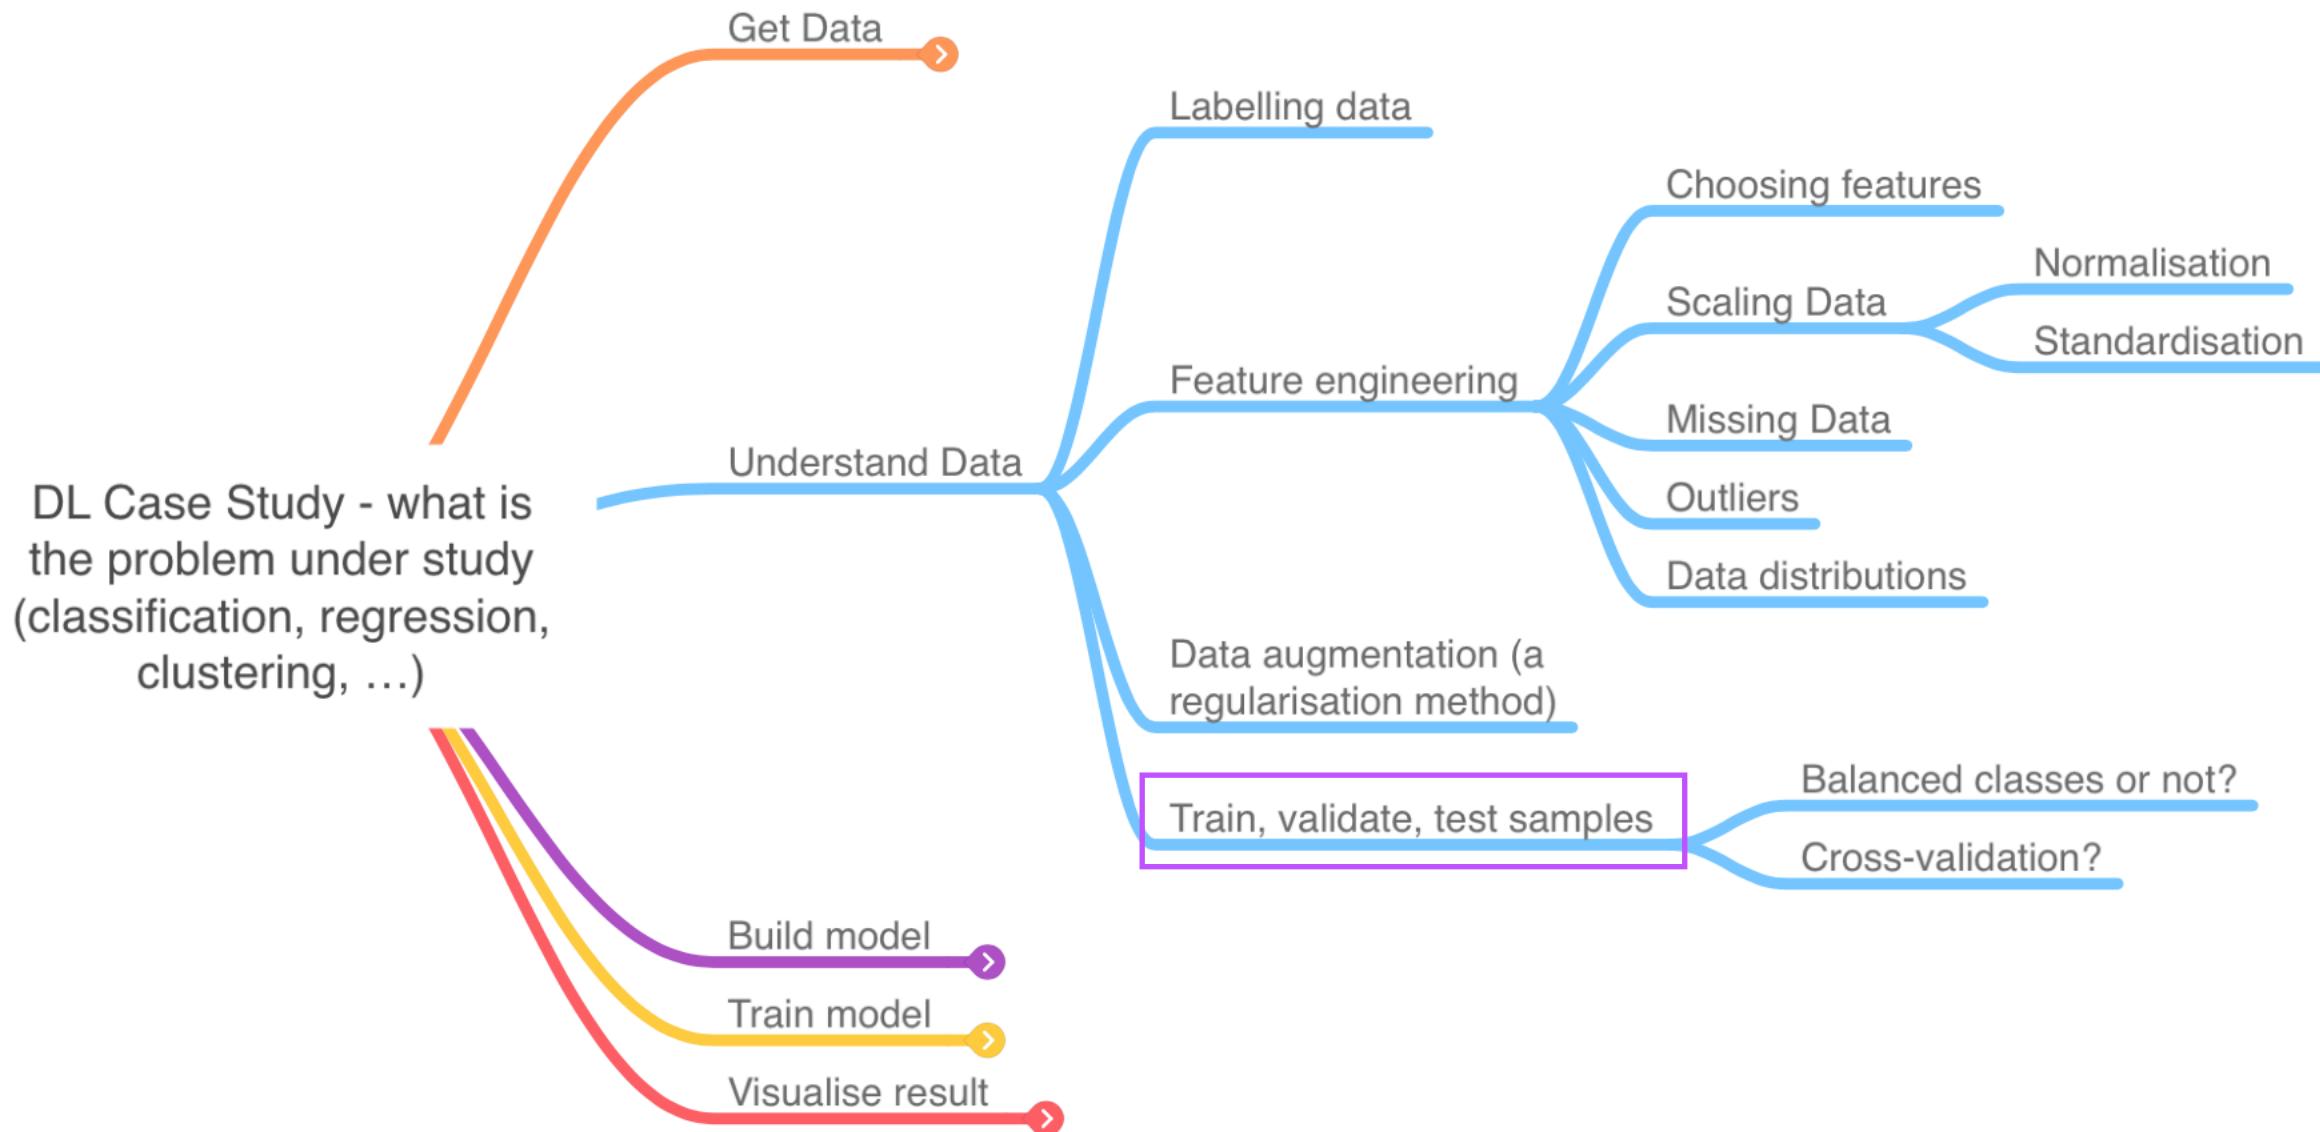
**Variance**: high variance means a lack of generalisation - the accuracy on a specific training or validation set may be high but the variance in accuracy across many sets is low



# Starting from the end

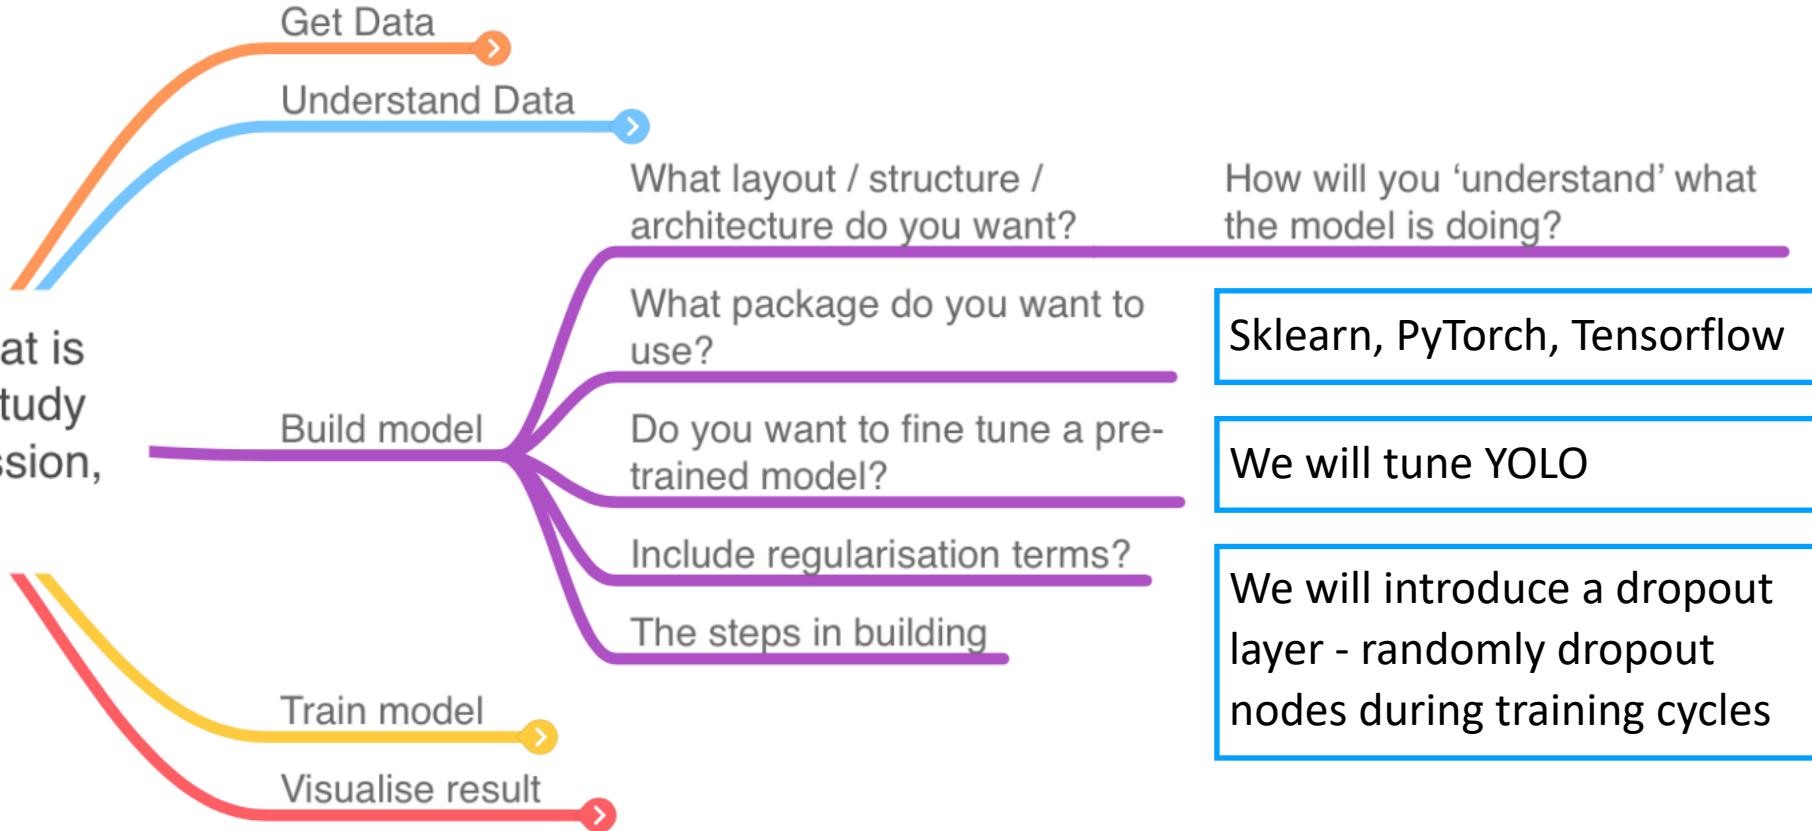


# Starting from the end



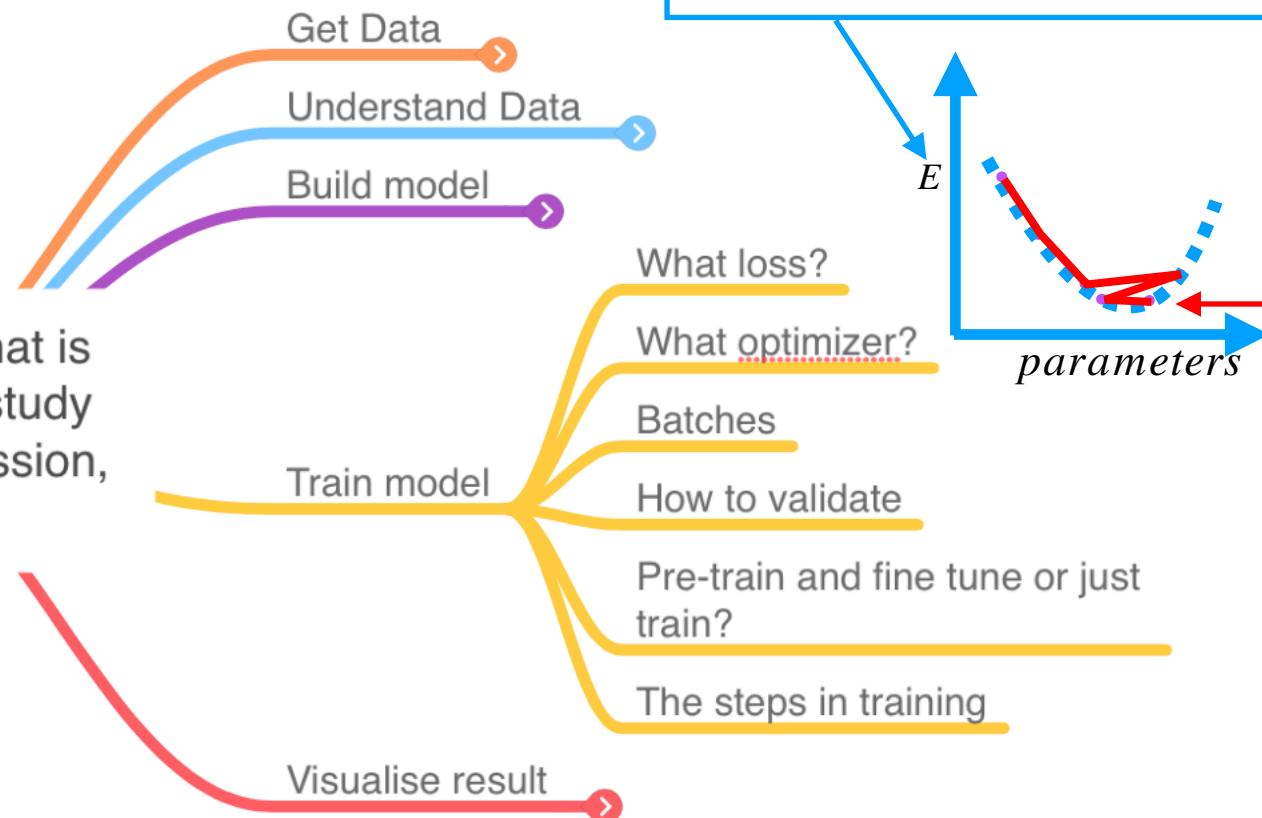
# Starting from the end

DL Case Study - what is the problem under study  
(classification, regression, clustering, ...)

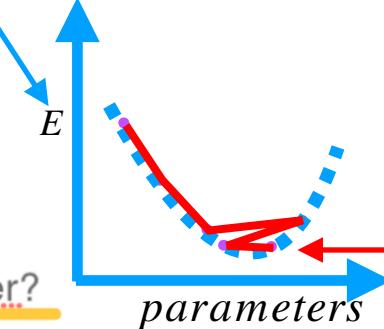


# Starting from the end

DL Case Study - what is the problem under study (classification, regression, clustering, ...)



Loss/cost - difference  
between your model outputs  
and the 'truth'



Our loss functions

We will use a cross entropy for classification (binary or categorical).

We will use mean squared error for regression.

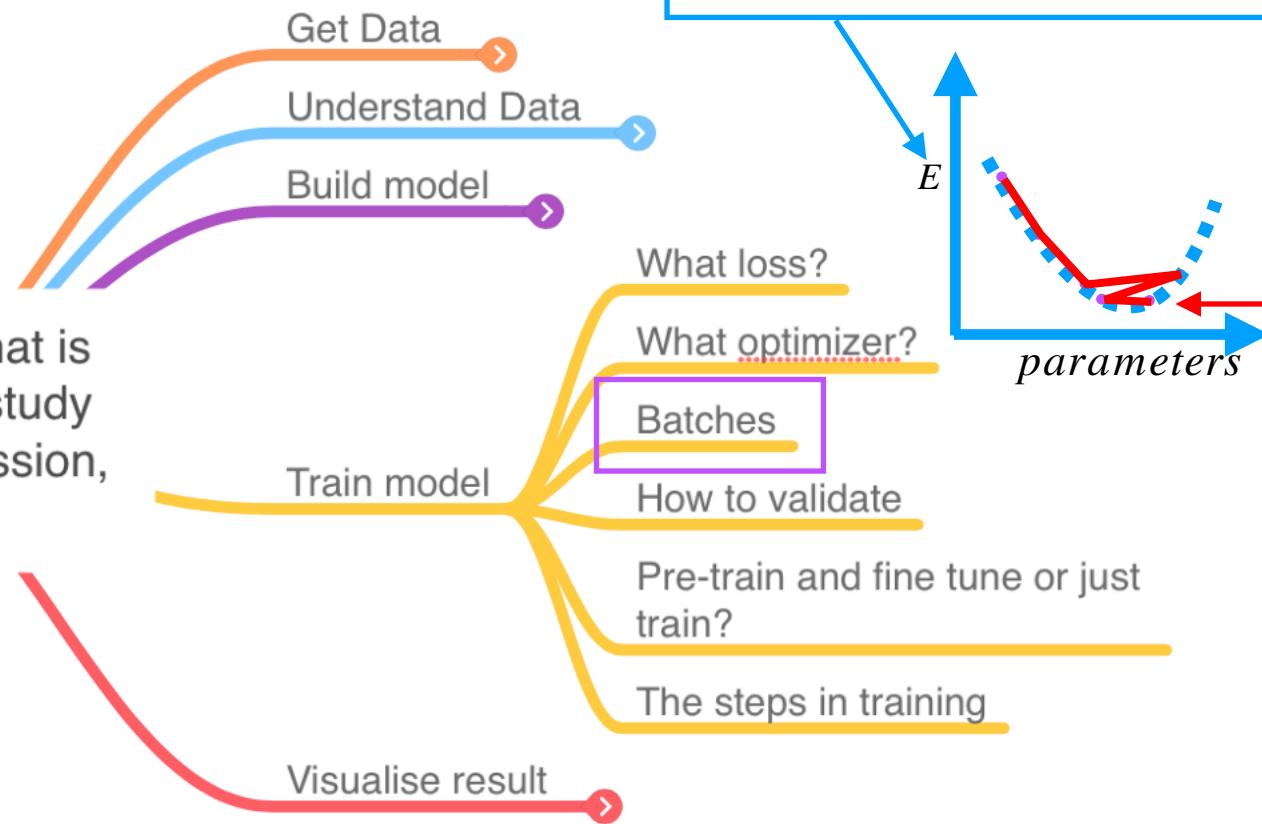
Optimisers tell us how to step about in our parameters space.

Typically gradient decent (you pick learning rate or 'step size') or adaptive (it adapts learning rate).

We will use Adam

# Starting from the end

DL Case Study - what is the problem under study (classification, regression, clustering, ...)



## Our loss functions

We will use a cross entropy for classification (binary or categorical).

We will use mean squared error for regression.

Optimisers tell us how to step about in our parameters space.

Typically gradient decent (you pick learning rate or 'step size') or adaptive (it adapts learning rate).

We will use Adam .

Generally different algorithms change how they process batches.

# Samples, batches, epochs

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

My input dataset

# Samples, batches, epochs

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \boxed{\text{My input dataset}} \quad [x_1] \quad \boxed{\text{A data sample}}$$

# Samples, batches, epochs

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

My input dataset

$$[x_1]$$

A data sample

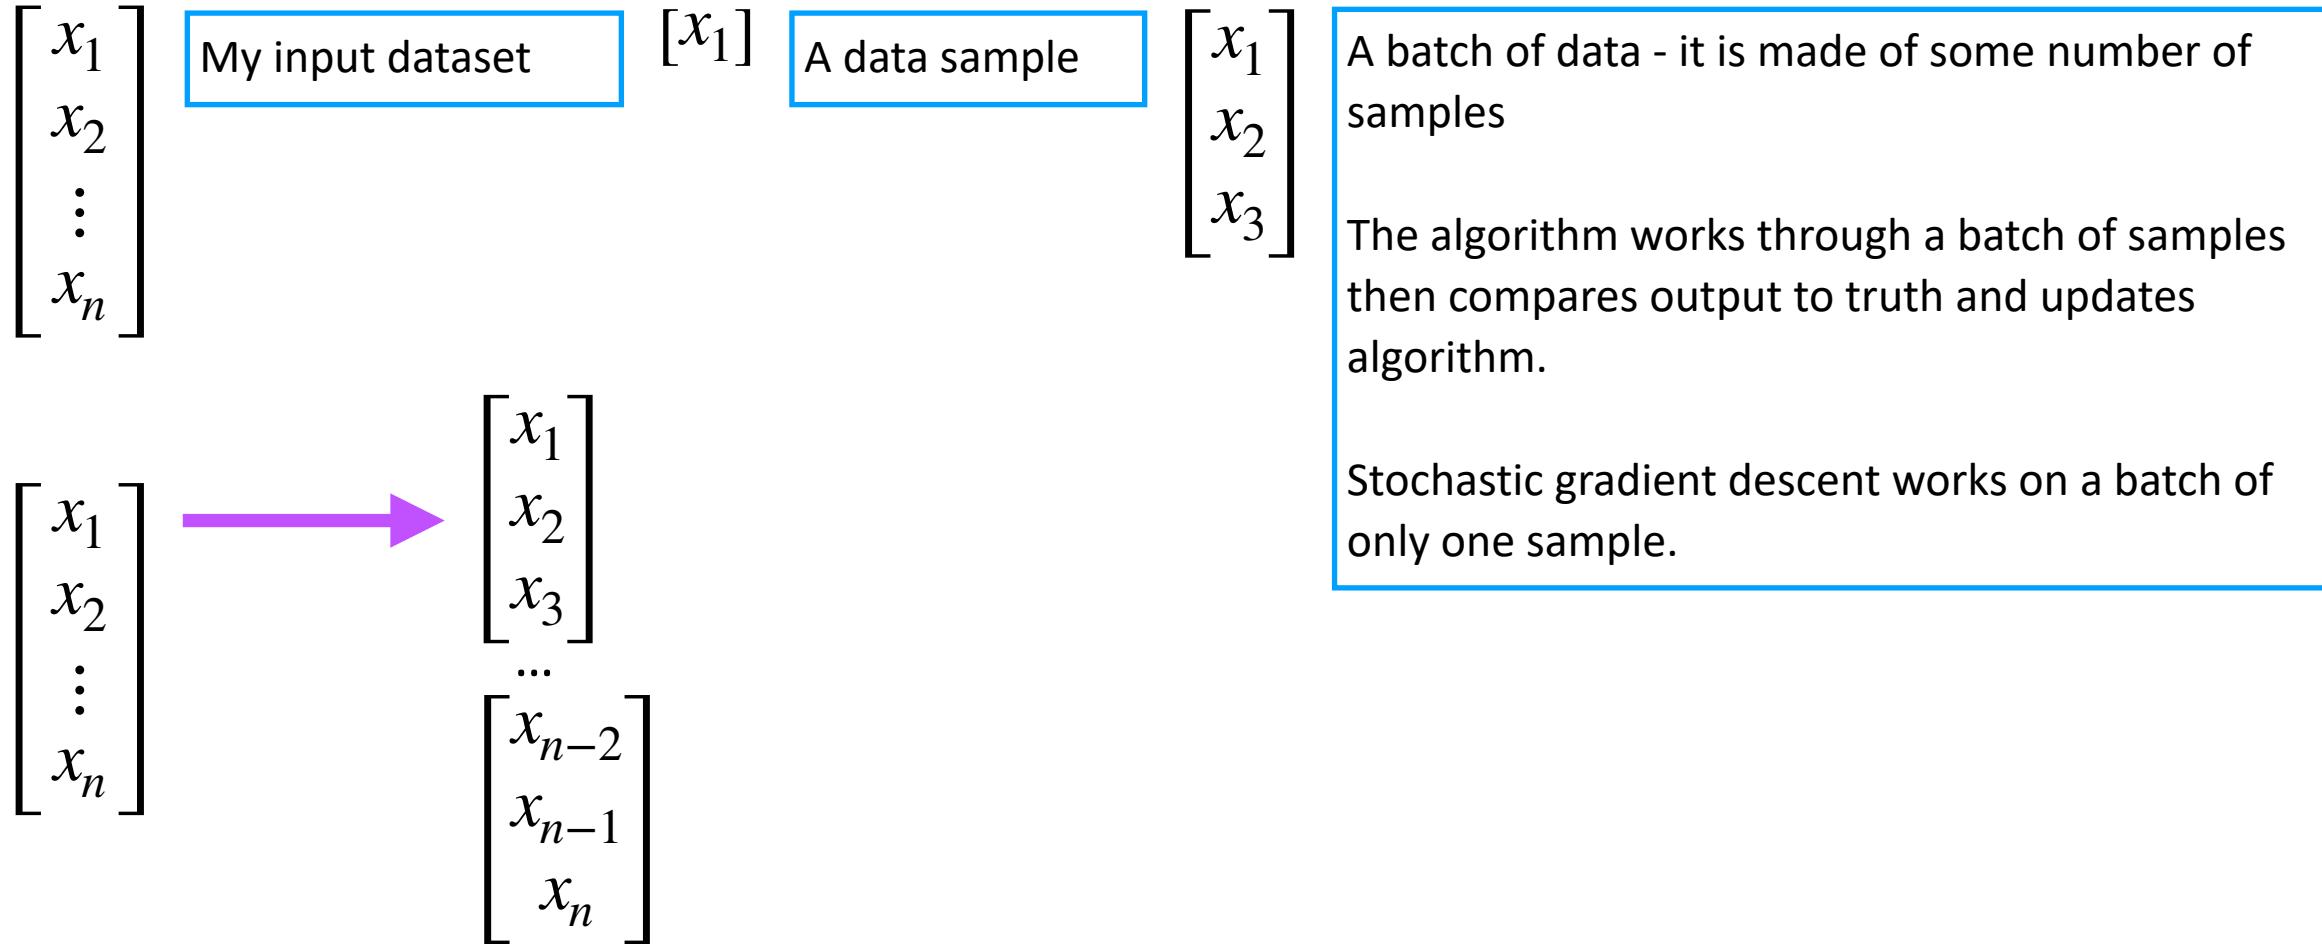
$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

A batch of data - it is made of some number of samples

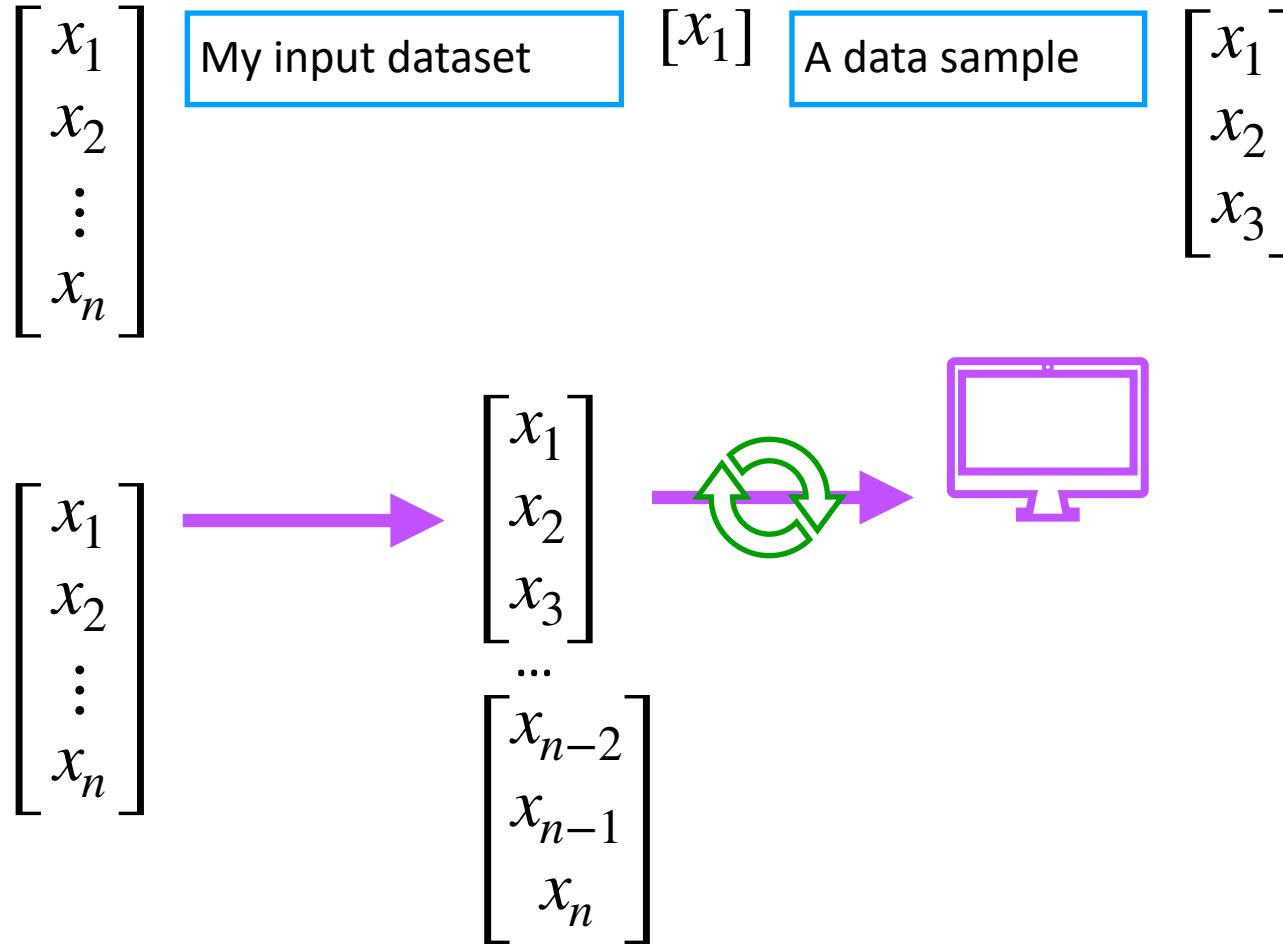
The algorithm works through a batch of samples then compares output to truth and updates algorithm.

Stochastic gradient descent works on a batch of only one sample.

# Samples, batches, epochs



# Samples, batches, epochs

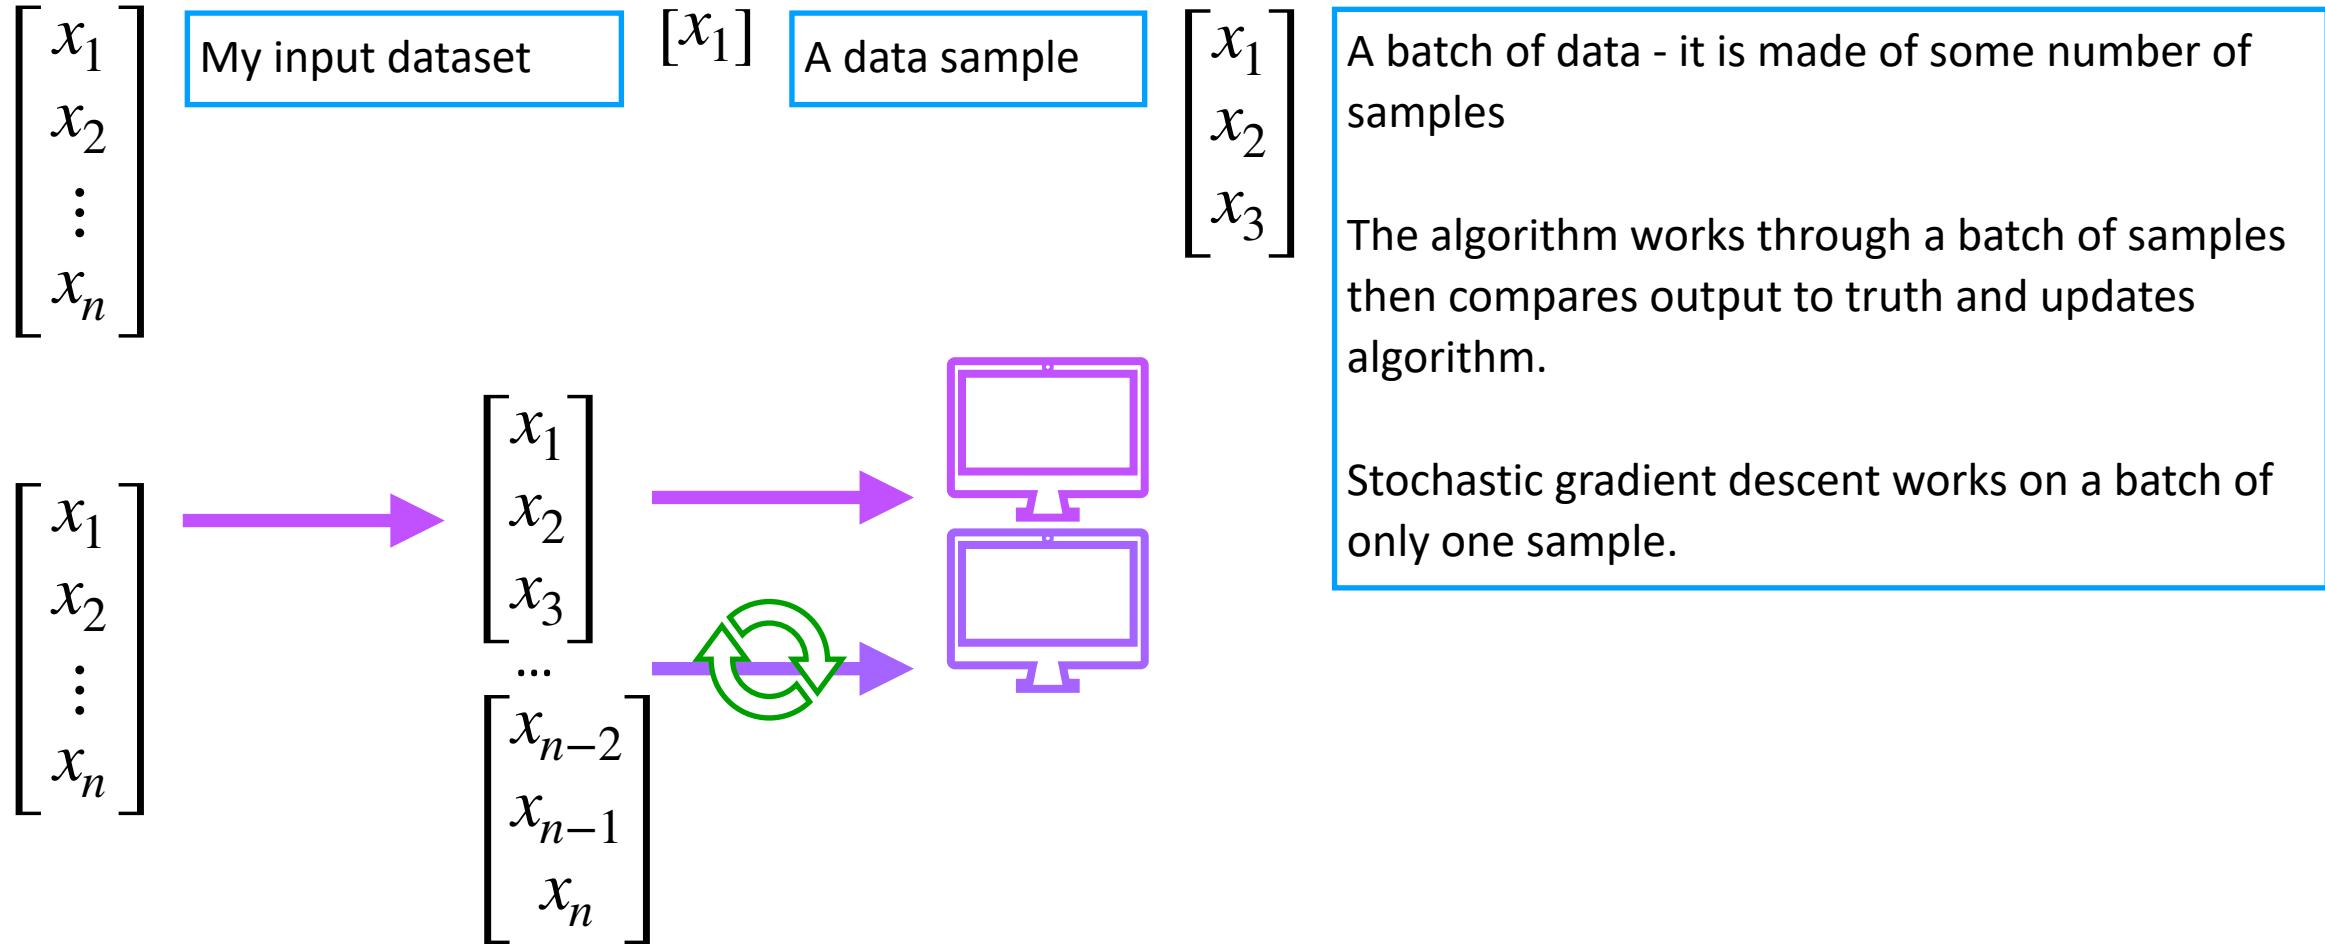


A batch of data - it is made of some number of samples

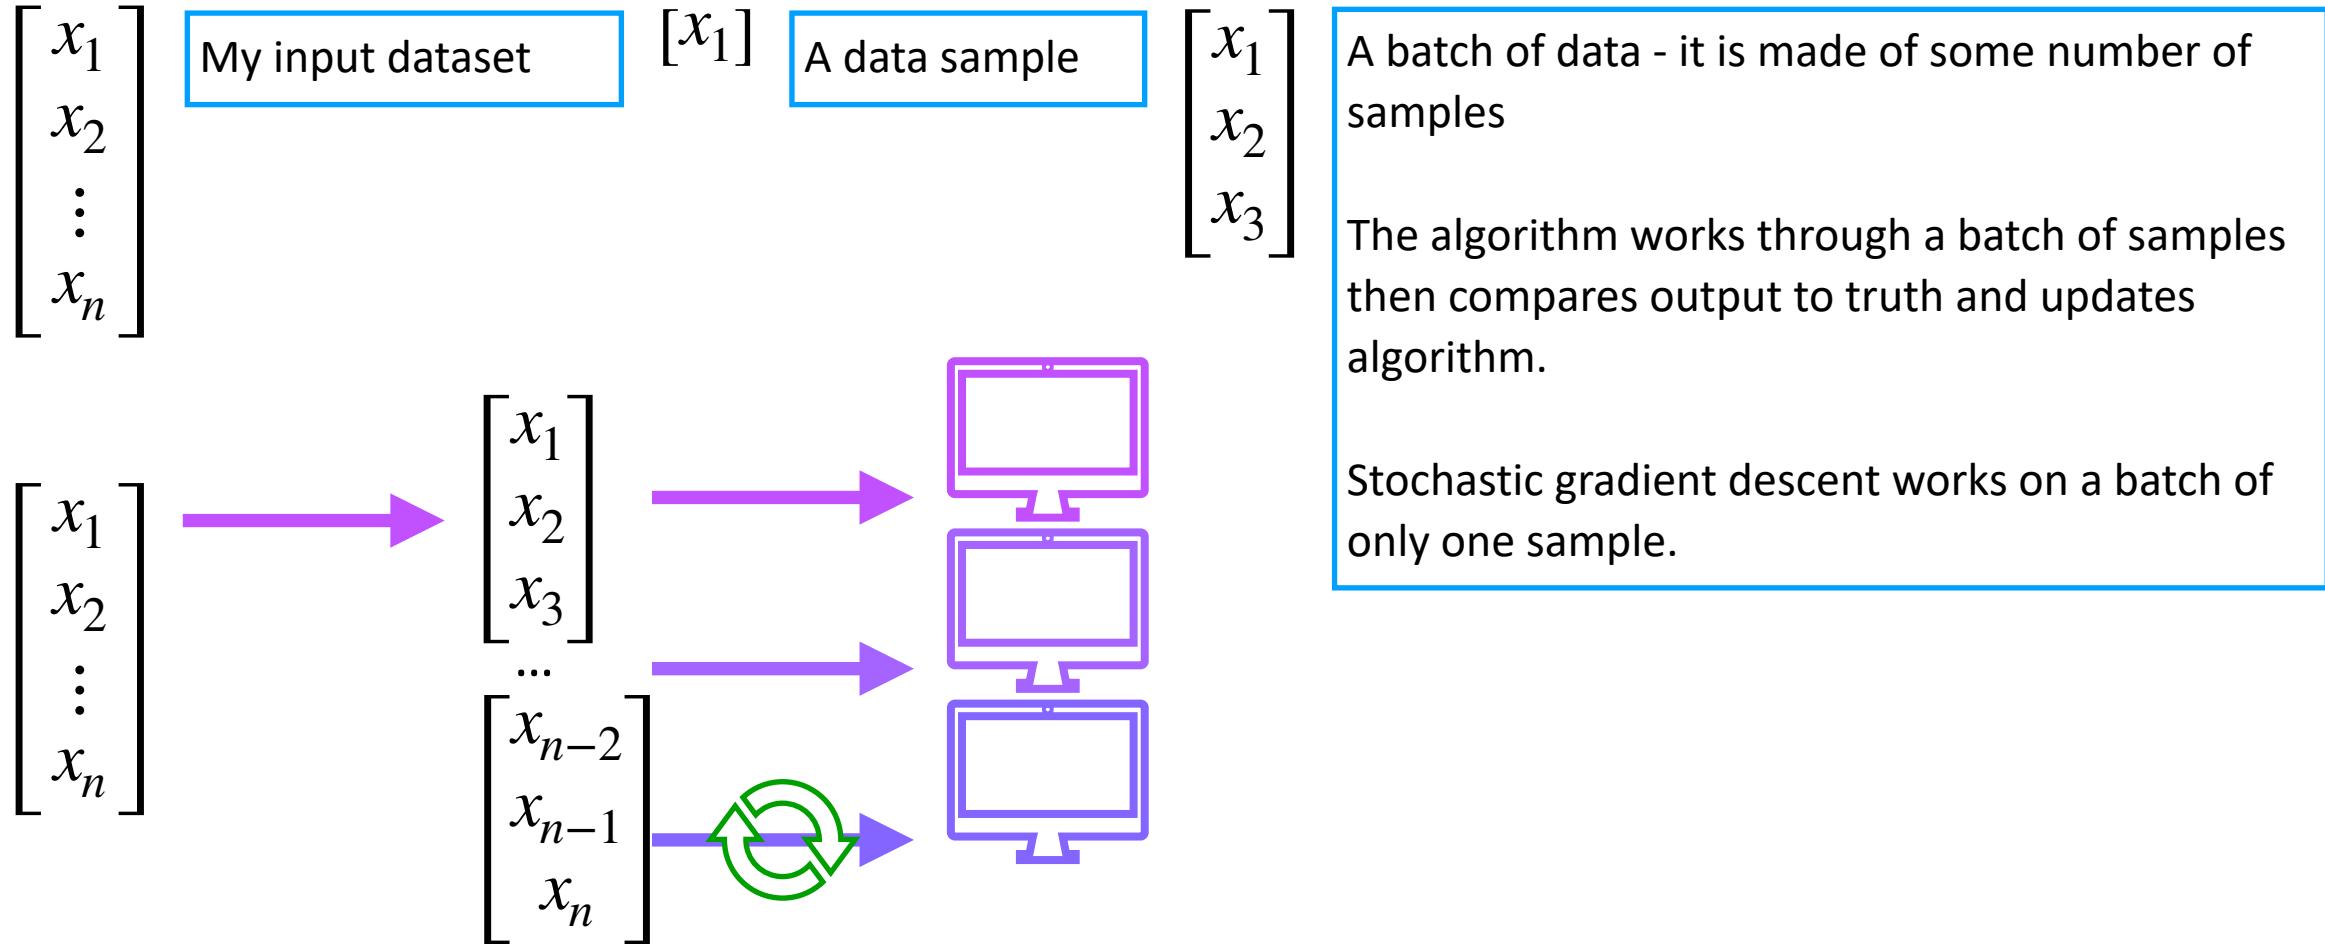
The algorithm works through a batch of samples then compares output to truth and updates algorithm.

Stochastic gradient descent works on a batch of only one sample.

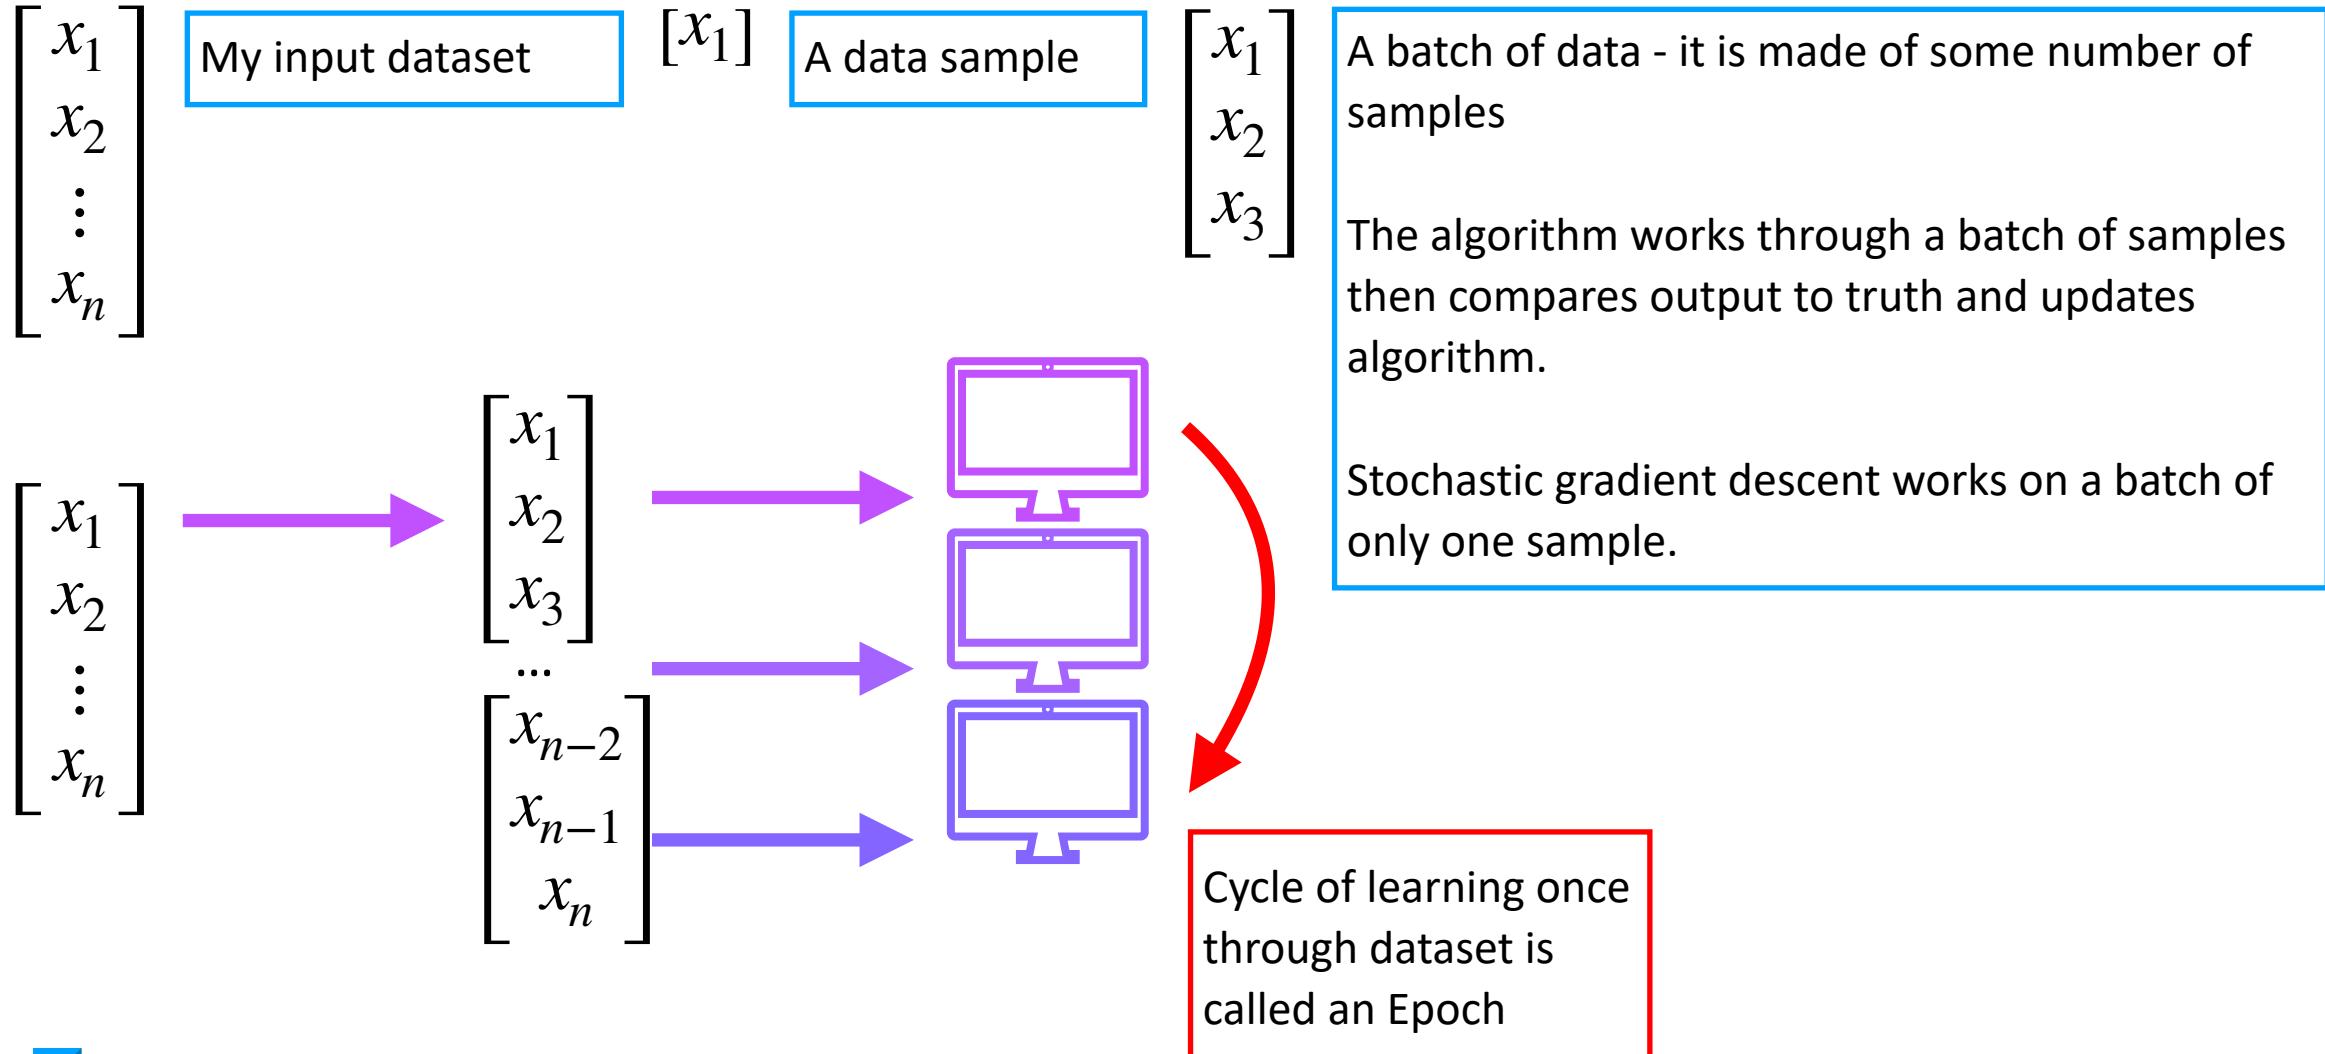
# Samples, batches, epochs



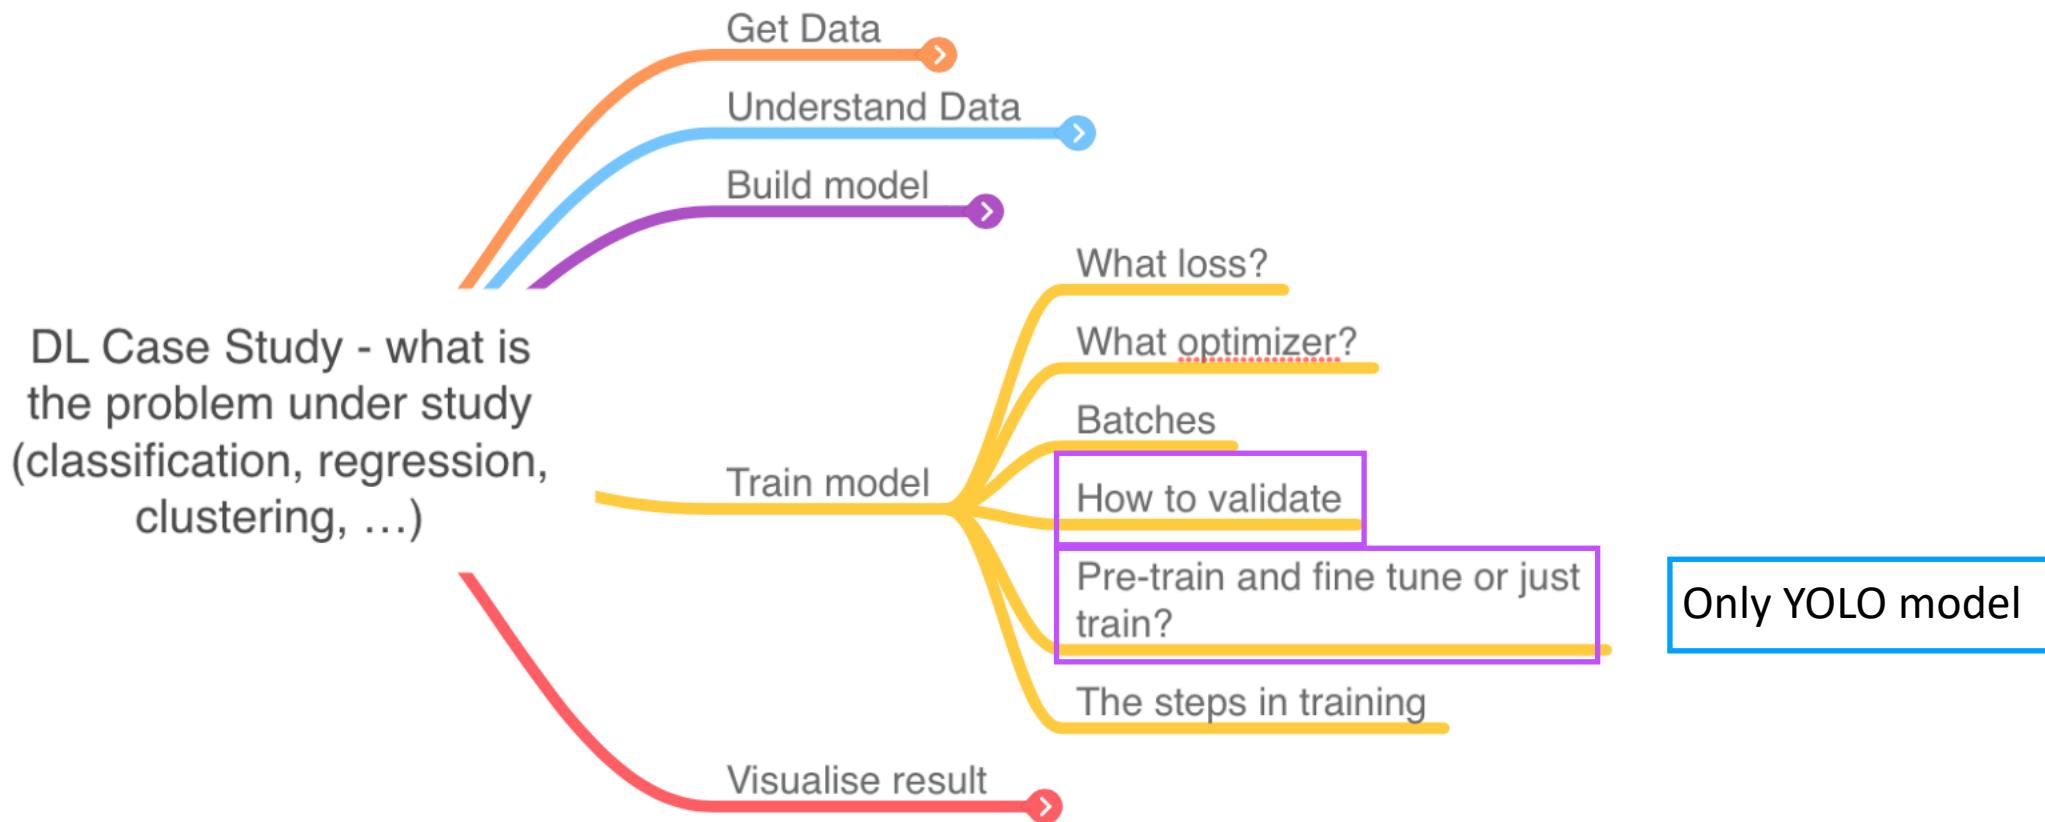
# Samples, batches, epochs



# Samples, batches, epochs

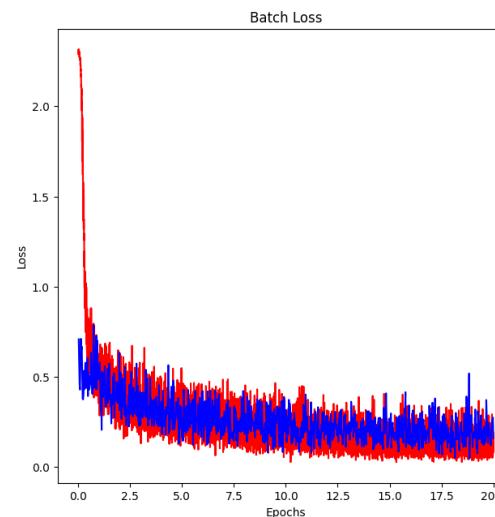
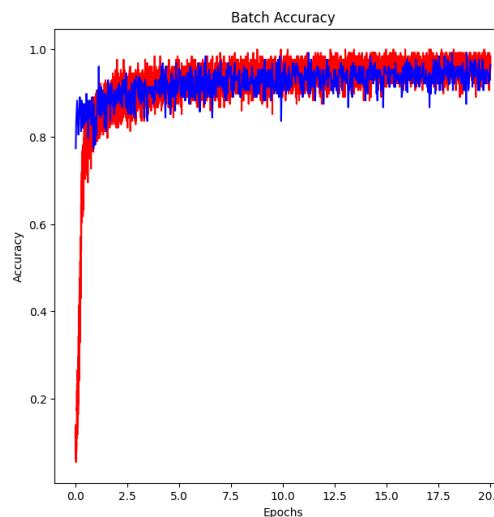
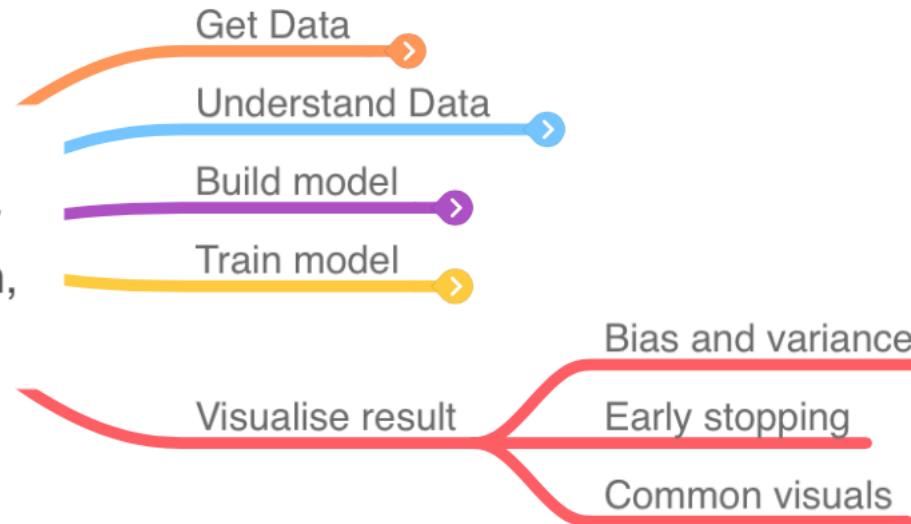


# Starting from the end

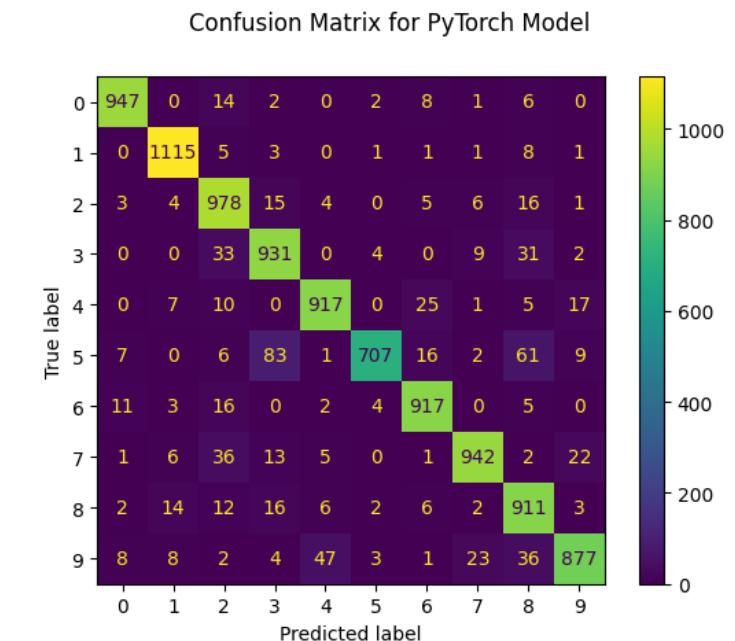


# Starting from the end

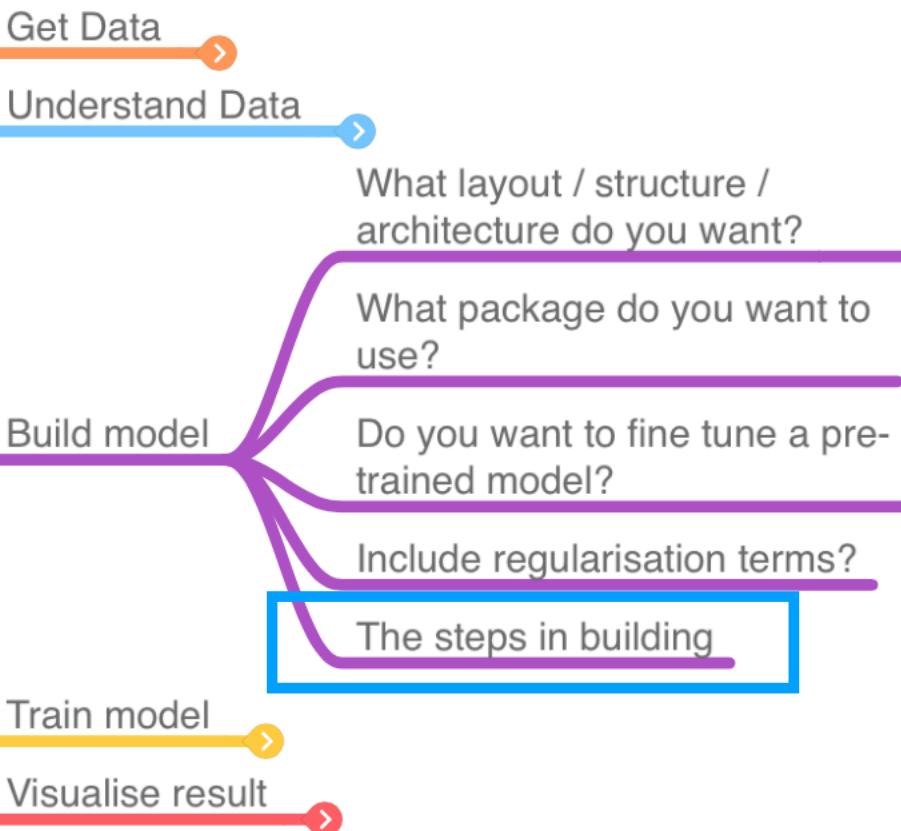
DL Case Study - what is the problem under study  
(classification, regression, clustering, ...)



There are packages for visualisation v popular one is tensor board for tensorflow - we won't have time to use these



# The steps in building



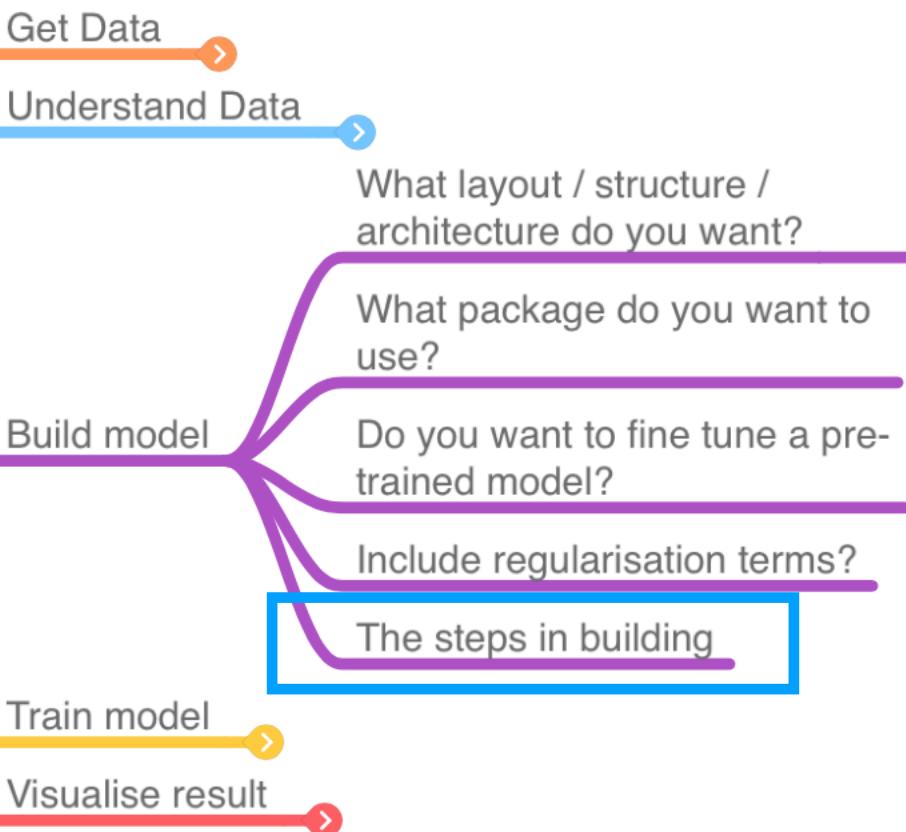
Objects - things with properties and methods

## PyTorch

- Make a 'class' for your model object
- A class consists of two parts:
  - `__init__()` - model components go here
  - `forward` - this part computes things

```
class TinyModel(torch.nn.Module):  
  
    def __init__(self):  
        super(TinyModel, self).__init__()  
  
        self.linear1 = torch.nn.Linear(100, 200)  
        self.activation = torch.nn.ReLU()  
        self.linear2 = torch.nn.Linear(200, 10)  
        self.softmax = torch.nn.Softmax()  
  
    def forward(self, x):  
        x = self.linear1(x)  
        x = self.activation(x)  
        x = self.linear2(x)  
        x = self.softmax(x)  
        return x
```

# The steps in building

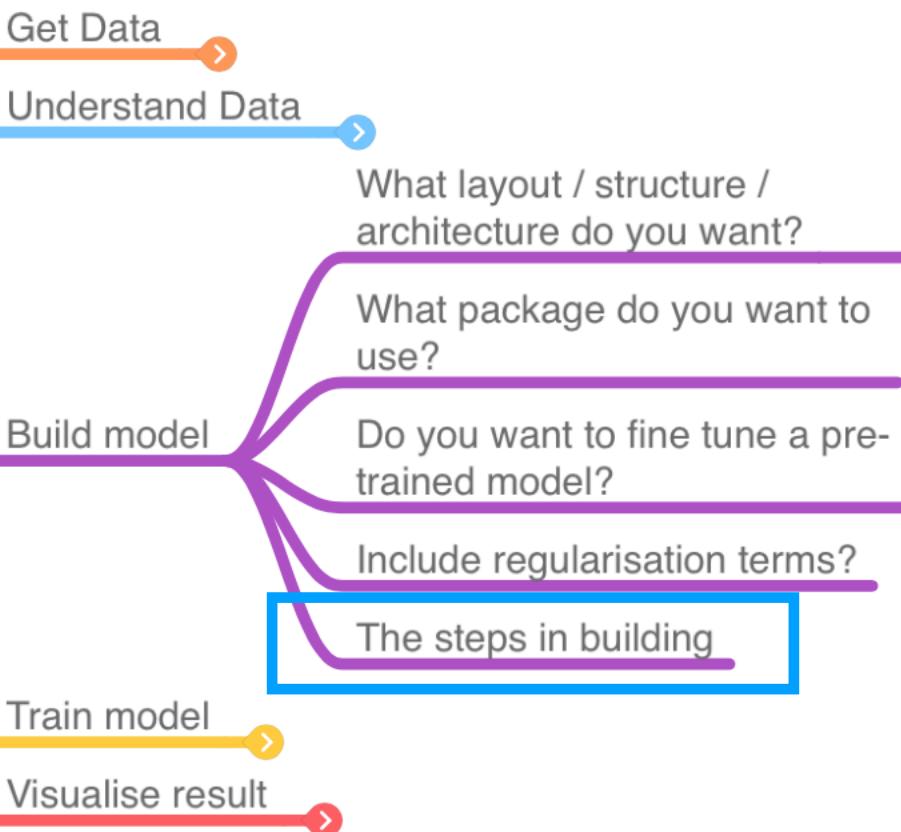


## PyTorch

- Make a 'class' for your model
- A class consists of two parts:
  - `__init__()` - model components go here
  - `forward` - this part computes things

```
class TinyModel(torch.nn.Module):  
  
    def __init__(self):  
        super(TinyModel, self).__init__()  
  
        self.linear1 = torch.nn.Linear(100, 200)  
        self.activation = torch.nn.ReLU()  
        self.linear2 = torch.nn.Linear(200, 10)  
        self.softmax = torch.nn.Softmax()  
  
    def forward(self, x):  
        x = self.linear1(x)  
        x = self.activation(x)  
        x = self.linear2(x)  
        x = self.softmax(x)  
        return x
```

# The steps in building



**Really good tutorial:**  
[https://pytorch.org/tutorials/beginner/introyt/modelsyt\\_tutorial.html](https://pytorch.org/tutorials/beginner/introyt/modelsyt_tutorial.html)

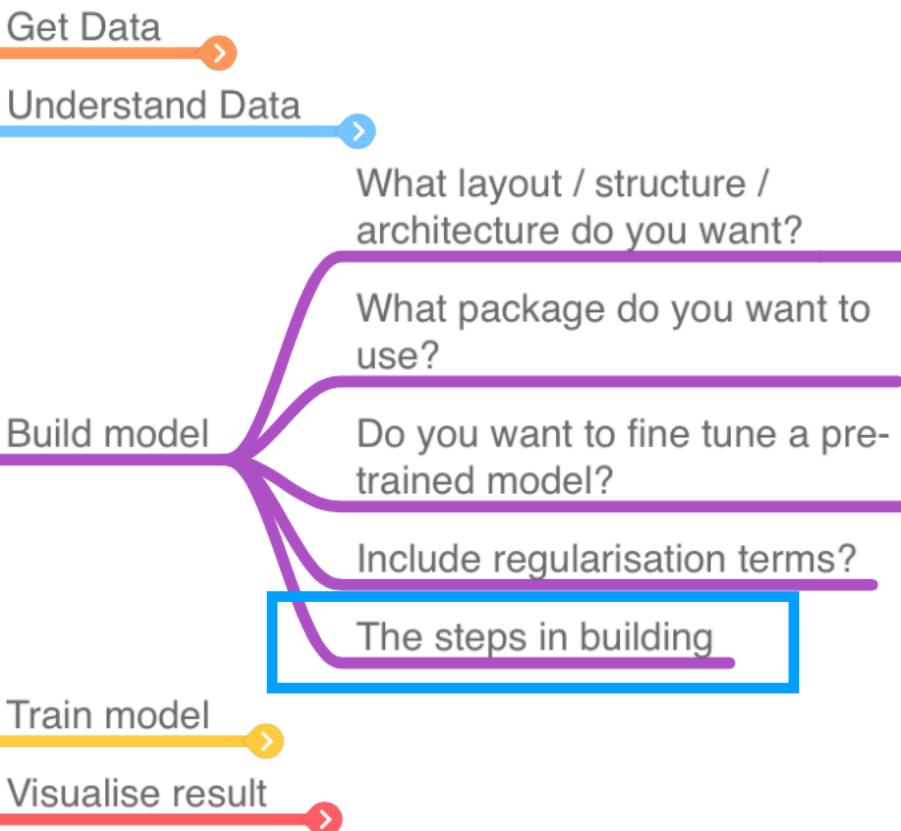
## PyTorch

- Make a 'class' for your model
- A class consists of two parts:
  - `__init__()` - model components go here
  - `forward` - this part computes things

```
class TinyModel(torch.nn.Module):  
  
    def __init__(self):  
        super(TinyModel, self).__init__()  
  
        self.linear1 = torch.nn.Linear(100, 200)  
        self.activation = torch.nn.ReLU()  
        self.linear2 = torch.nn.Linear(200, 10)  
        self.softmax = torch.nn.Softmax()  
  
    def forward(self, x):  
        x = self.linear1(x)  
        x = self.activation(x)  
        x = self.linear2(x)  
        x = self.softmax(x)  
        return x
```

Sometimes seen as:  
`super().__init__()`  
Gives the  
`torch.nn.Module` class  
ability to call this function

# The steps in building



```
import torch.functional as F

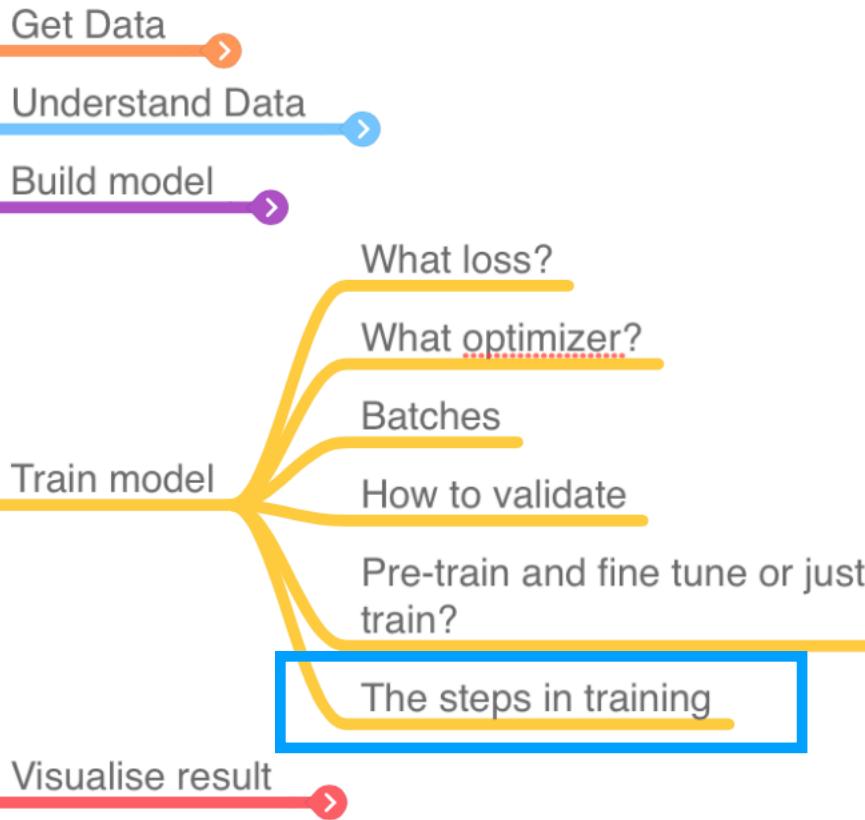
class LeNet(torch.nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        # 1 input image channel (black & white), 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = torch.nn.Conv2d(1, 6, 5)
        self.conv2 = torch.nn.Conv2d(6, 16, 3)
        # an affine operation: y = Wx + b
        self.fc1 = torch.nn.Linear(16 * 6 * 6, 120) # 6*6 from image dimension
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

# The steps in training

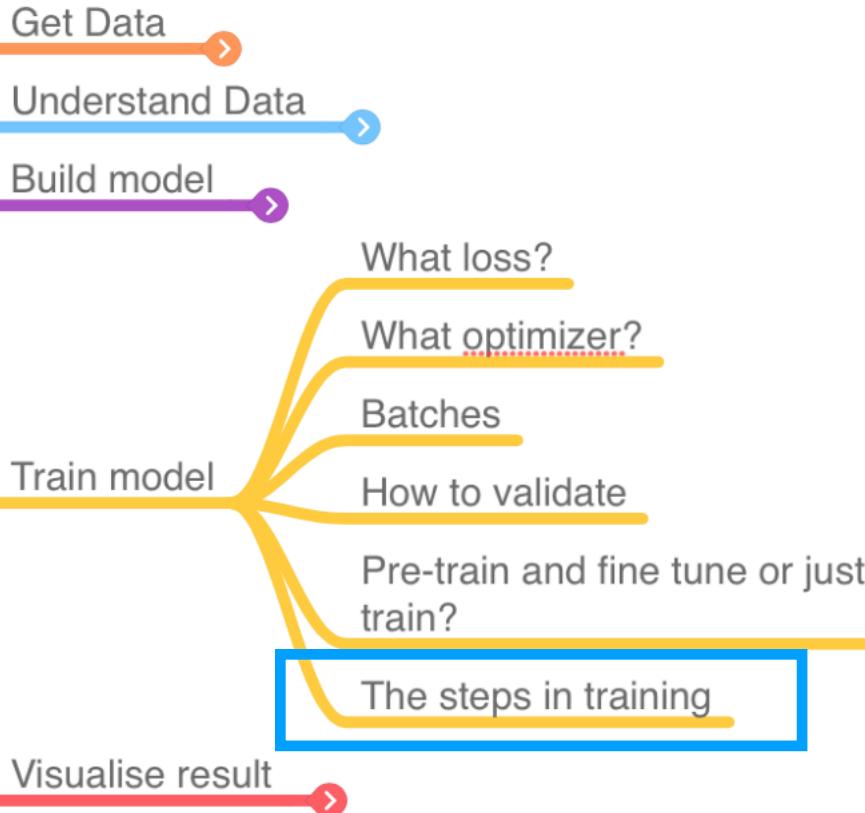


- Put model into 'train' mode: `model.train()`
- For each 'batch':
  - Reset gradient to zero - incase not
  - Predict outputs from inputs
  - Calculate loss
  - Compute gradients using backpropagation
  - Update parameters
  - Compute accuracy and loss

**Really good doc to read:**

<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html>

# The steps in training



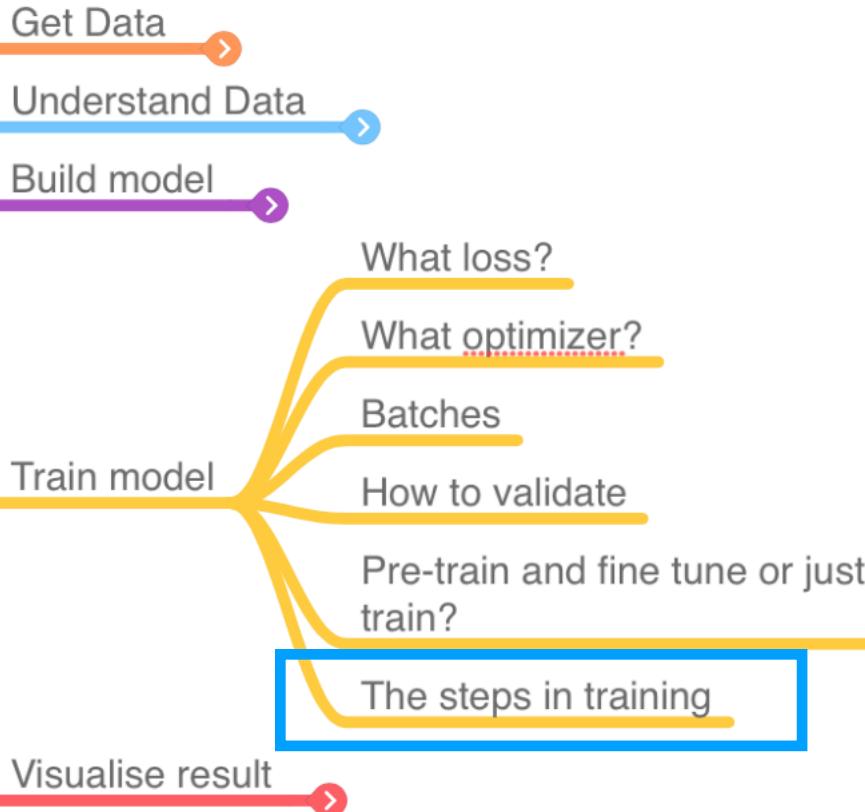
- Put model into 'train' mode: `model.train()`
- For each 'batch':
  - Reset gradient to zero - incase not
  - Predict outputs from inputs
  - Calculate loss
  - Compute gradients using backpropagation
  - Update parameters
  - Compute accuracy and loss

```
model.train() # model is now in 'train' mode
for epoch in range(5):
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad() #reset gradient to zero
        outputs = model(images) # predict outputs from inputs
        loss = nn.CrossEntropyLoss()(outputs, labels) # calculate loss
        loss.backward() # compute gradients using backpropagation
        optimizer.step() # update parameters
```

**Really good doc to read:**

<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html>

# The steps in training



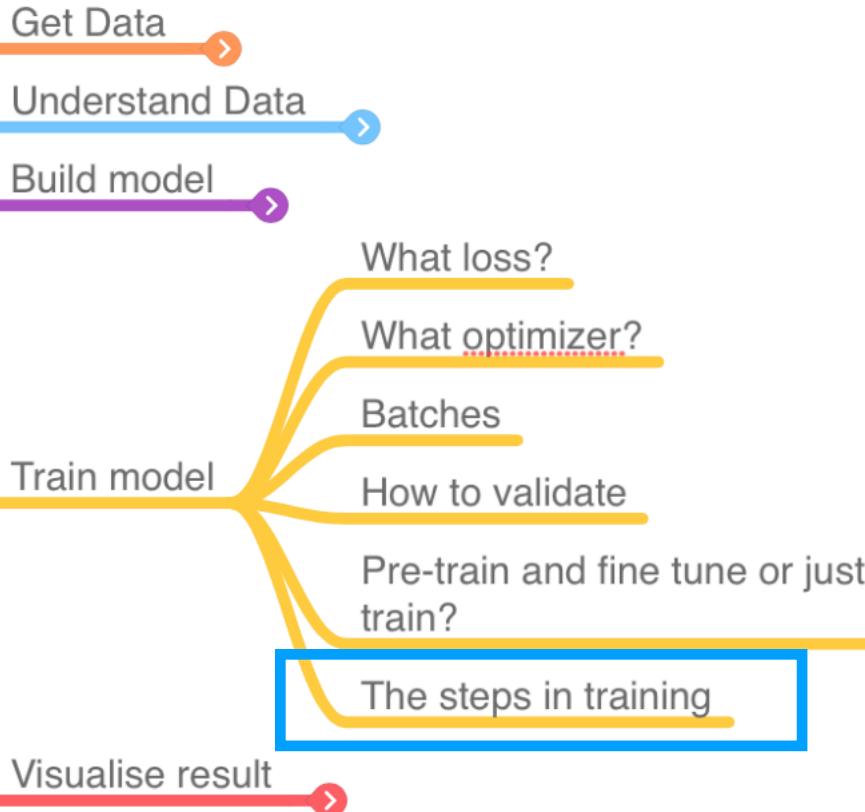
- Put model into 'train' mode: `model.train()`
- For each 'batch':
  - Reset gradient to zero - incase not
  - Predict outputs from inputs
  - Calculate loss
  - Compute gradients using backpropagation
  - Update parameters
  - Compute accuracy and loss

```
model.train() # model is now in 'train' mode
for epoch in range(5):
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad() #reset gradient to zero
        outputs = model(images) # predict outputs from inputs
        loss = nn.CrossEntropyLoss()(outputs, labels) # calculate loss
        loss.backward() # compute gradients using backpropagation
        optimizer.step() # update parameters
```

**Really good doc to read:**

<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html>

# The steps in training



- Put model into 'train' mode: `model.train()`
- For each 'batch':
  - Reset gradient to zero - incase not
  - Predict outputs from inputs
  - Calculate loss
  - Compute gradients using backpropagation
  - Update parameters
  - Compute accuracy and loss

```
model.train() # model is now in 'train' mode
for epoch in range(5):
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad() #reset gradient to zero
        outputs = model(images) # predict outputs from inputs
        loss = nn.CrossEntropyLoss()(outputs, labels) # calculate loss
        loss.backward() # compute gradients using backpropagation
        optimizer.step() # update parameters
```

**Really good doc to read:**

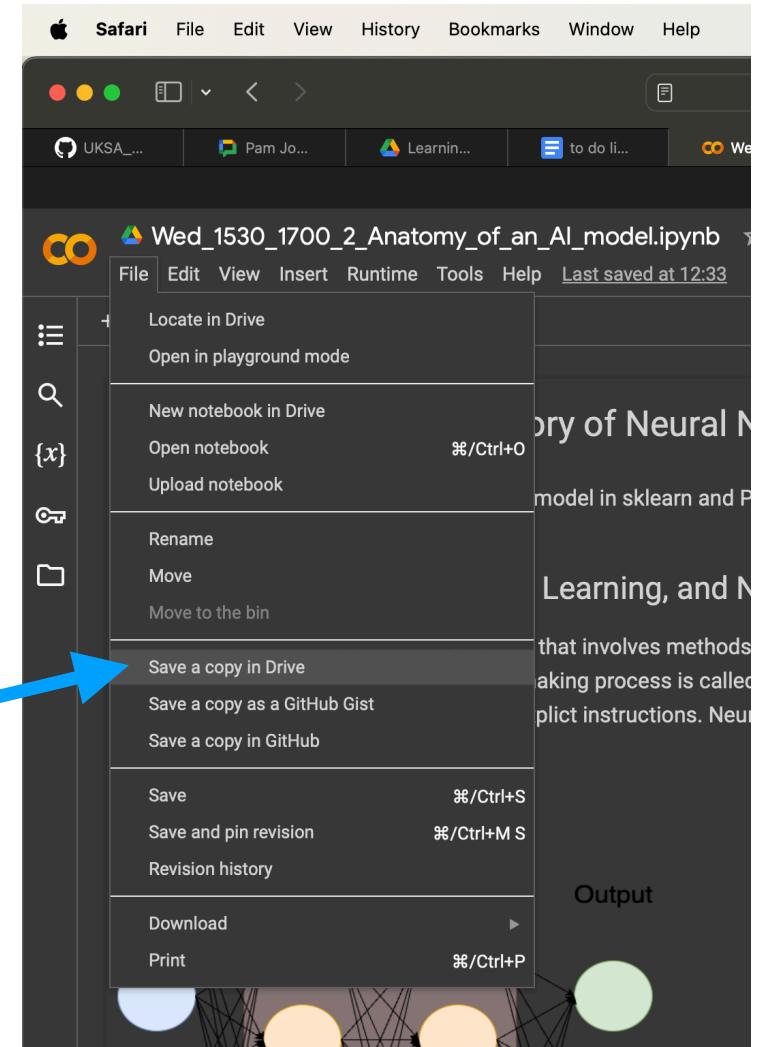
<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html>

# AI 101 - Wed + Thurs

- **TASK:** Build a model in PyTorch from scratch to classify the fashionMNIST dataset
  - **Aim to understand and apply:**
    - Exploring the data - how many categories and labels - why do we do this - how might data distributions affect your model output?
    - Anatomy of a model - steps for a model, steps for training
    - Key steps to evaluate model progress
  - **For reference or for the aficionados!**
    - 2 Clustering notebooks - sklearn and NN (autoencoders)
  - **DL - Thurs 11am - an example applied to remote sensing data**
    - **Aim:** to understand the jargon so you could pull this tutorial apart and understand its key features

# Plan A (ish)

- If on windows or linux: <https://148.197.234.100:32444/> sign-up on Jupyter
- If on Mac... Username and password, ssh port forwarding e.g. I would do...
  - ssh -L 8443:localhost:32444 becky@148.197.234.100



# Plan B

- Go to Moodle -> Go to Phase 2 -> Click link for Plan B  
(or click this link [https://drive.google.com/drive/folders/1InS4kCUwKO7gAI\\_TUjEcADtDSifb\\_aq8b?usp=share\\_link](https://drive.google.com/drive/folders/1InS4kCUwKO7gAI_TUjEcADtDSifb_aq8b?usp=share_link))
- Double click on a notebook: it will open in colab
- Go File -> Save a copy in Drive: this will open a new copy on your browser - edit this copy as you please